



Go Memory Management In Depth.

Billing&Rating DevTalks

- 1.Previously
- 2.Physical Memory
- 3.TCMalloc
- 4.Go Memory Structures
- 5.Garbage Collectors
- 6.Garbage Collector In Go
- 7.Tips



Calling convention

- All arguments passed on the stack
 - Offsets from FP
- Return arguments follow input arguments
 - Start of return arguments aligned to pointer size
- All registers are caller saved, except:
 - Stack pointer register
 - Zero register
 - G context pointer register
 - Frame pointer

Benchmark

```
//go:noinline
func ByRef(name string) *Account {
    acc := Account{
        name:     name,
        address: "some account address here",
        details: "some account details here",
        id:       12312312,
        zipCode:  143010,
    }

    return &acc
}
```

```
//go:noinline
func ByValue(name string) Account {
    return Account{
        name:     name,
        address: "some account address here",
        details: "some account details here",
        id:       12312312,
        zipCode:  143010,
    }
}
```

```
⇒ go test -v -bench=. -benchmem
```

```
goos: darwin
goarch: amd64
pkg: TeckTalk3/cmd/ref_value
cpu: Intel(R) Core(TM) i9-8950HK CPU @ 2.90GHz
```

```
BenchmarkByRef
```

Benchmark	Operations	Time	Memory	Allocations
BenchmarkByRef-12	26352846	39.25 ns/op	64 B/op	1 allocs/op

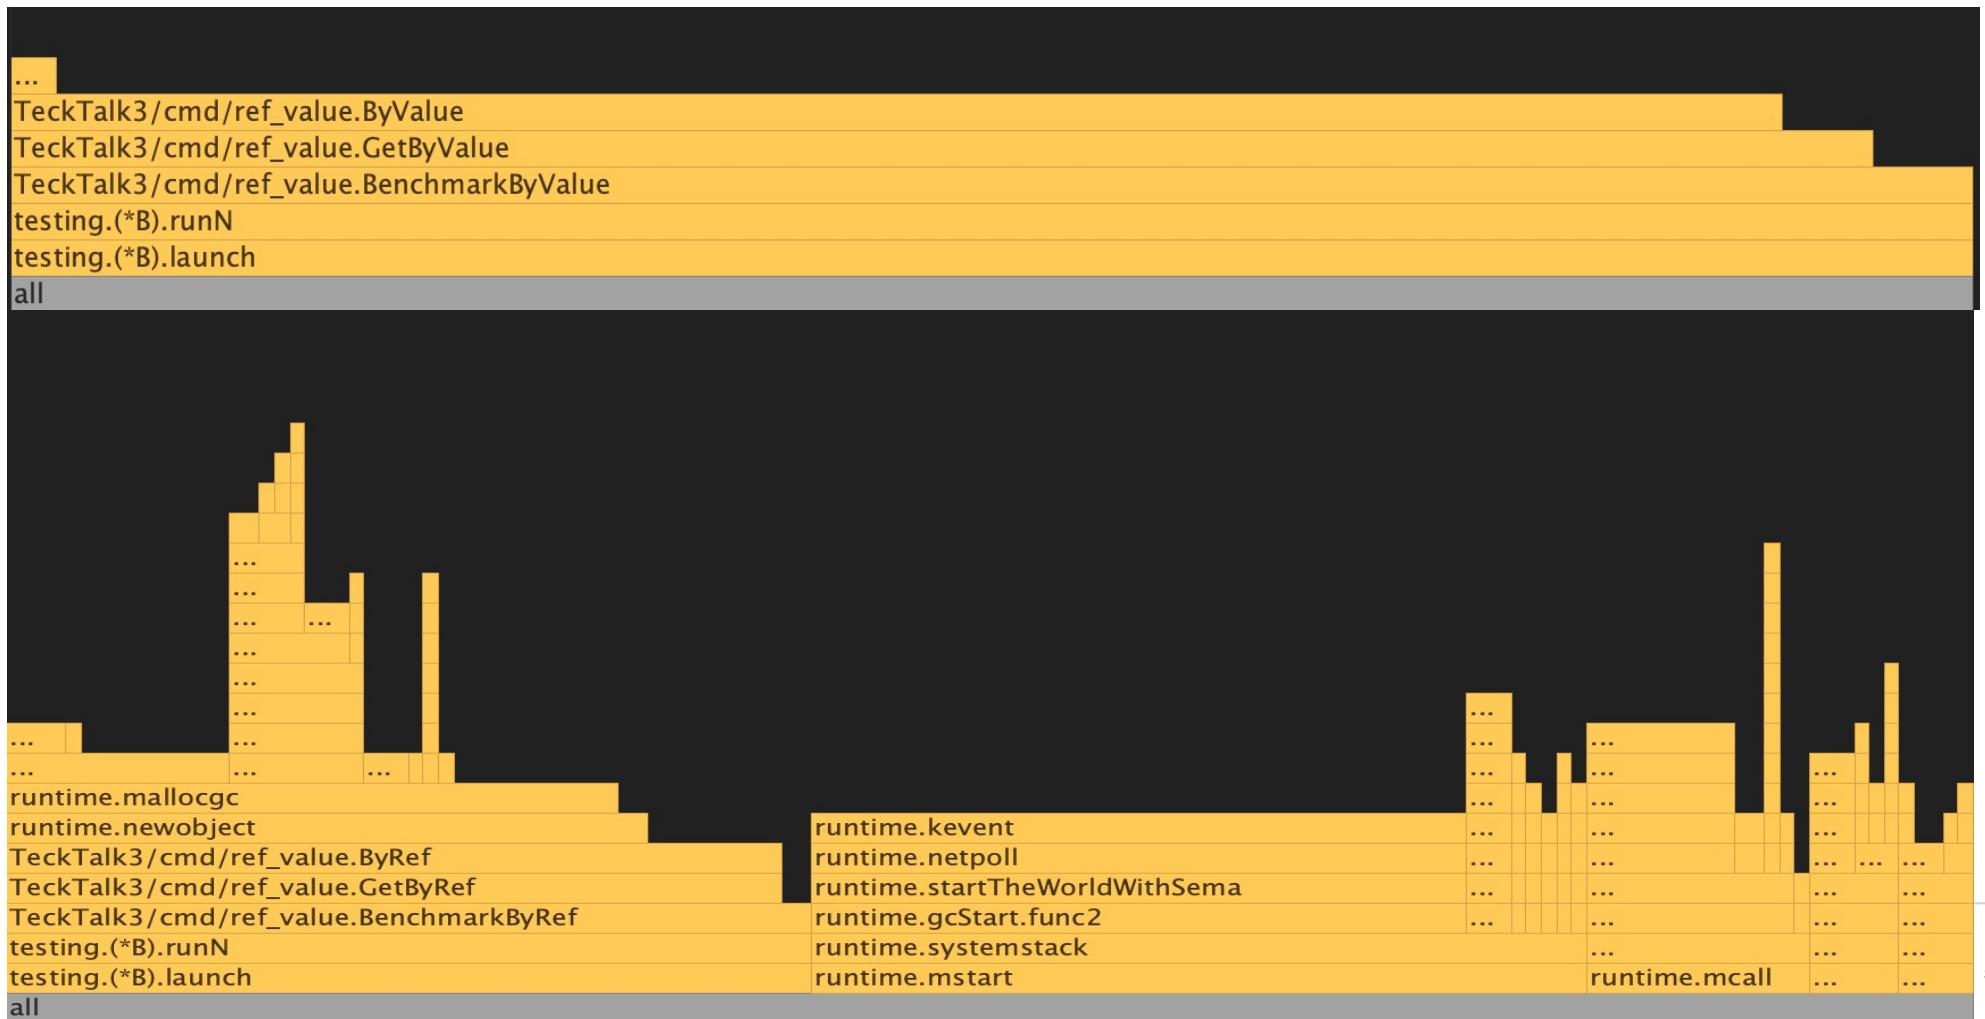
```
BenchmarkByValue
```

BenchmarkByValue-12	154288635	7.856 ns/op	0 B/op	0 allocs/op
---------------------	-----------	-------------	--------	-------------

```
PASS
```

```
ok      TeckTalk3/cmd/ref_value 3.375s
```

CPU Profile



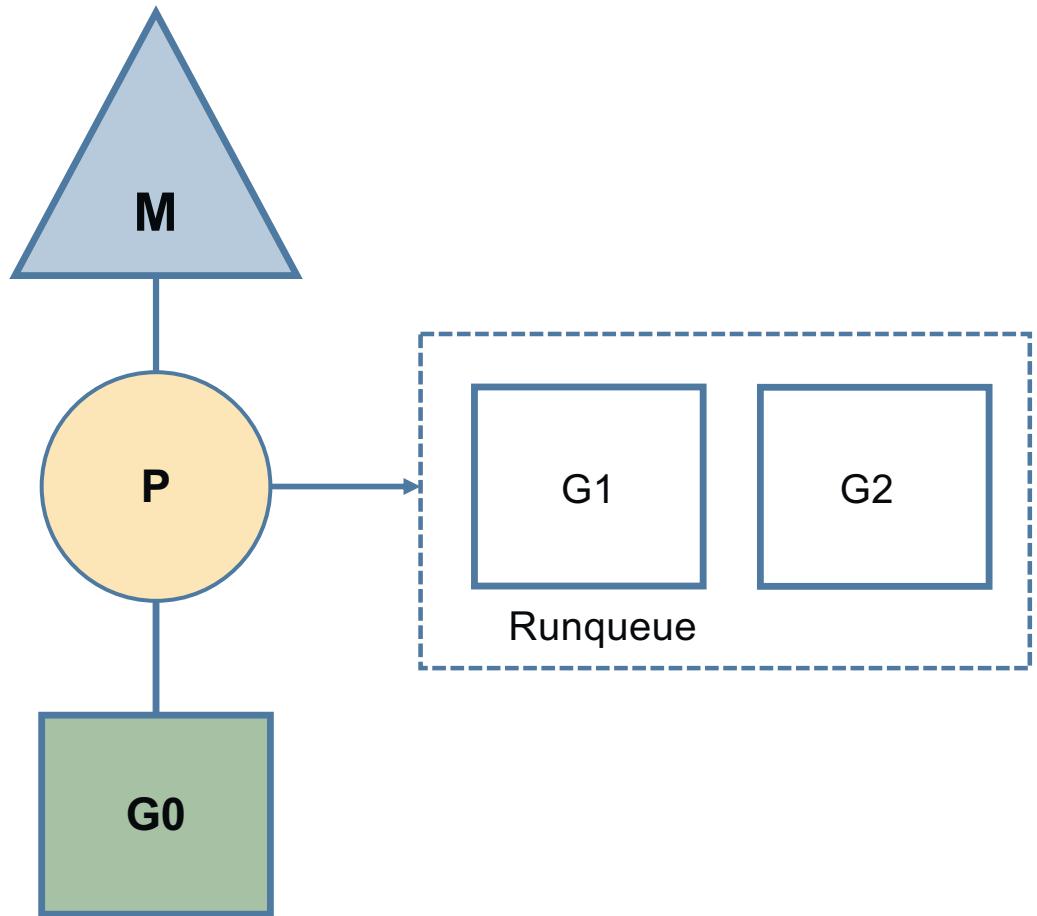


Concurrency

- Multiple OS threads
- Multiple goroutines
- User interact only with gorounites

G M P

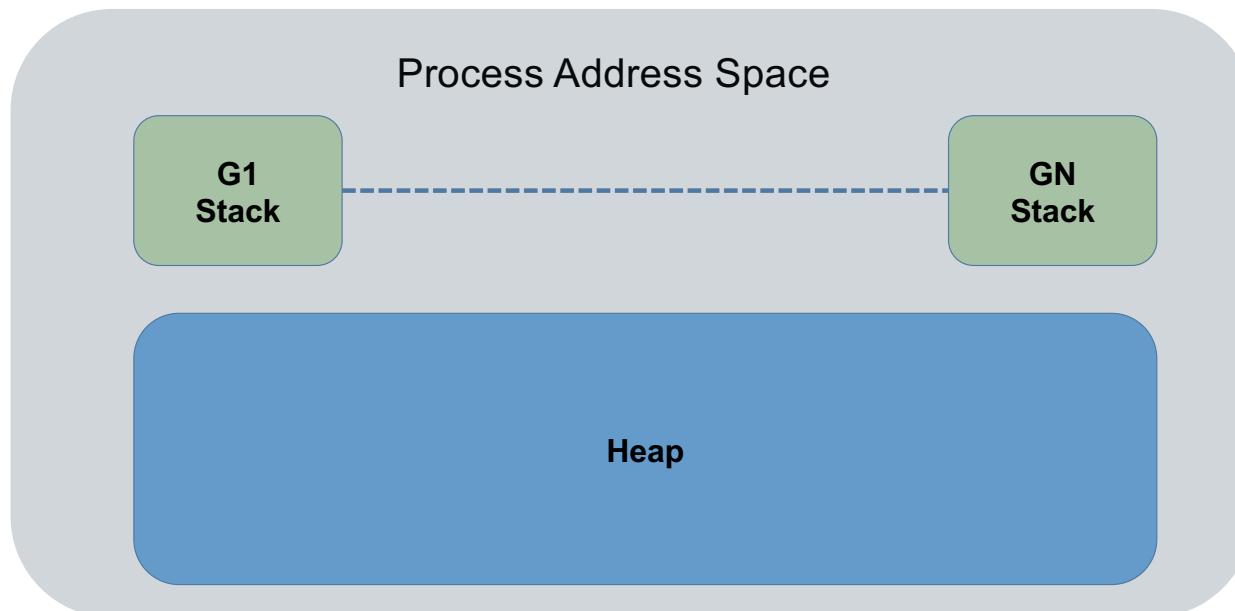
- M (machine) – OS thread
- P (processor) – Execution Context, holds runqueue
- G – Goroutine





The Go stack and heap

- Threads managed by the OS. Go use goroutines.
- Goroutines exist within user space - runtime, and not the OS, sets the rules of how stacks behave.



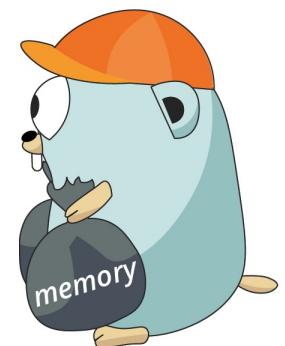
Magic Inside

```
//go:noinline
func ByRef(name string) *Account {
    acc := Account{
        name:     name,
        address: "some account address here",
        details: "some account details here",
        id:       12312312,
        zipCode: 143010,
    }

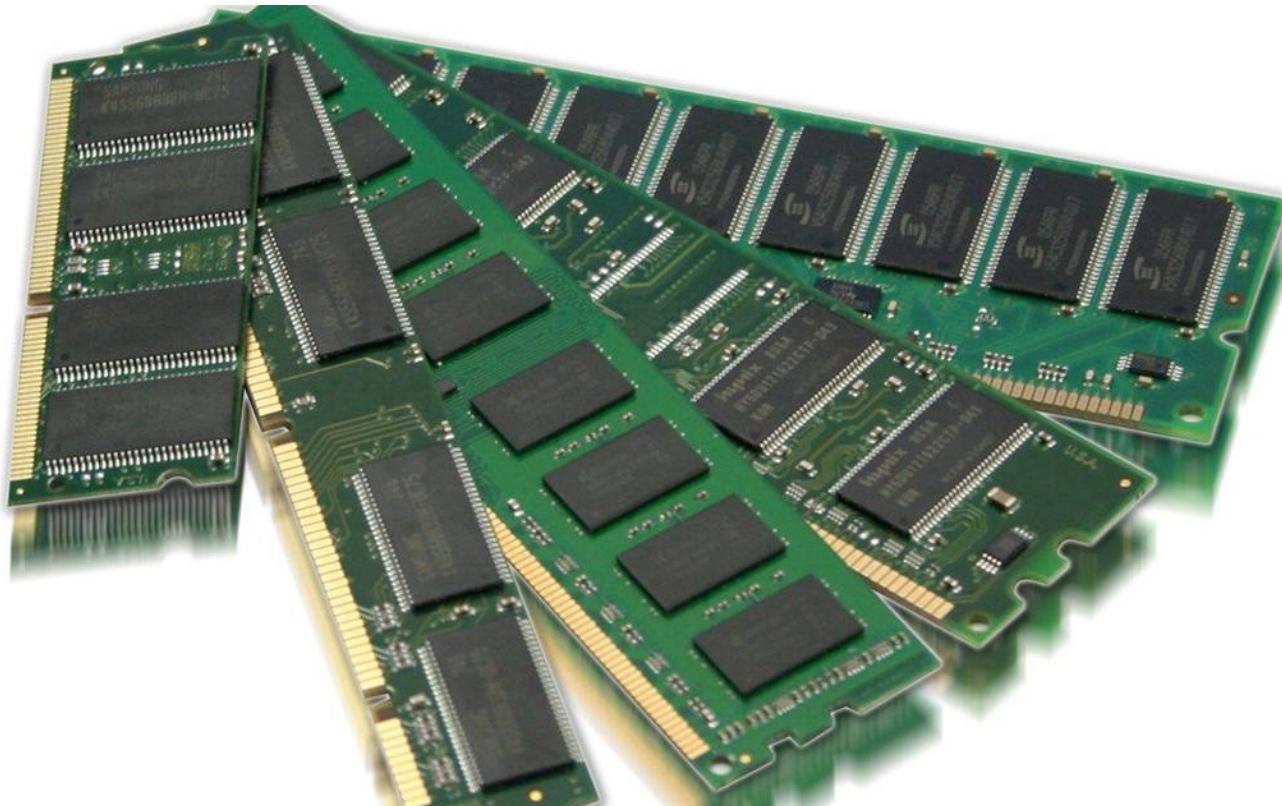
    return &acc
}
```



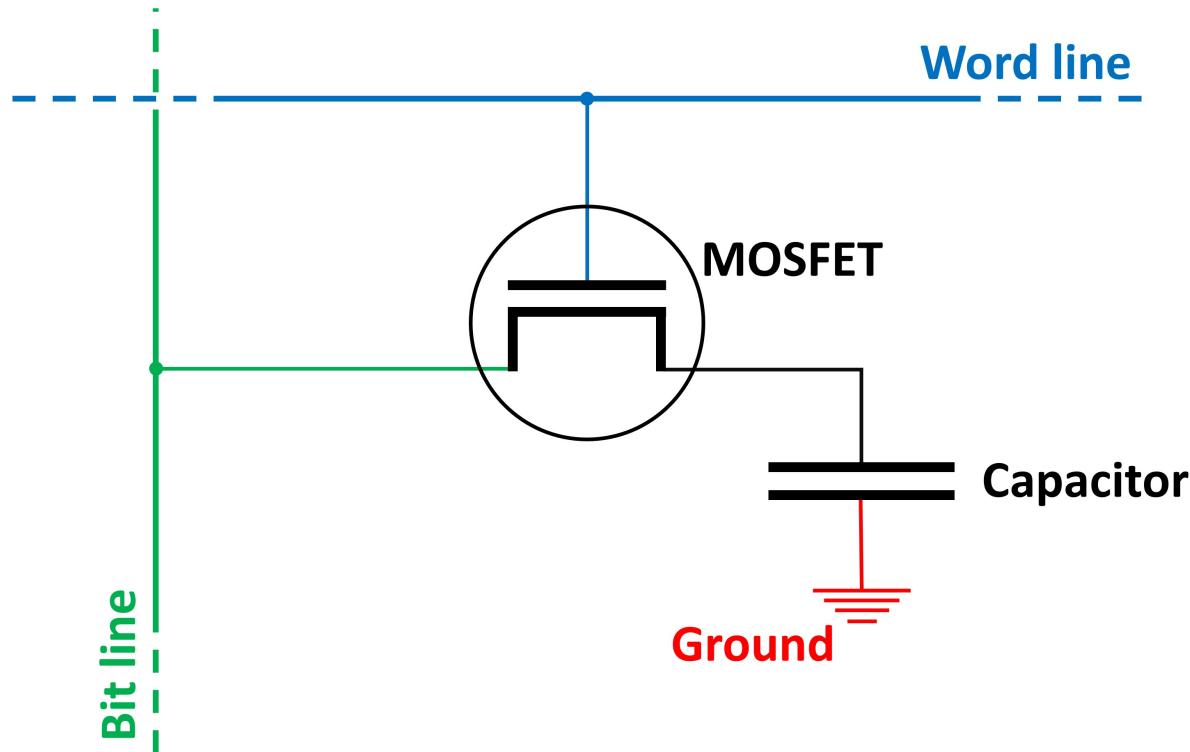
go magic()



Physical Memory

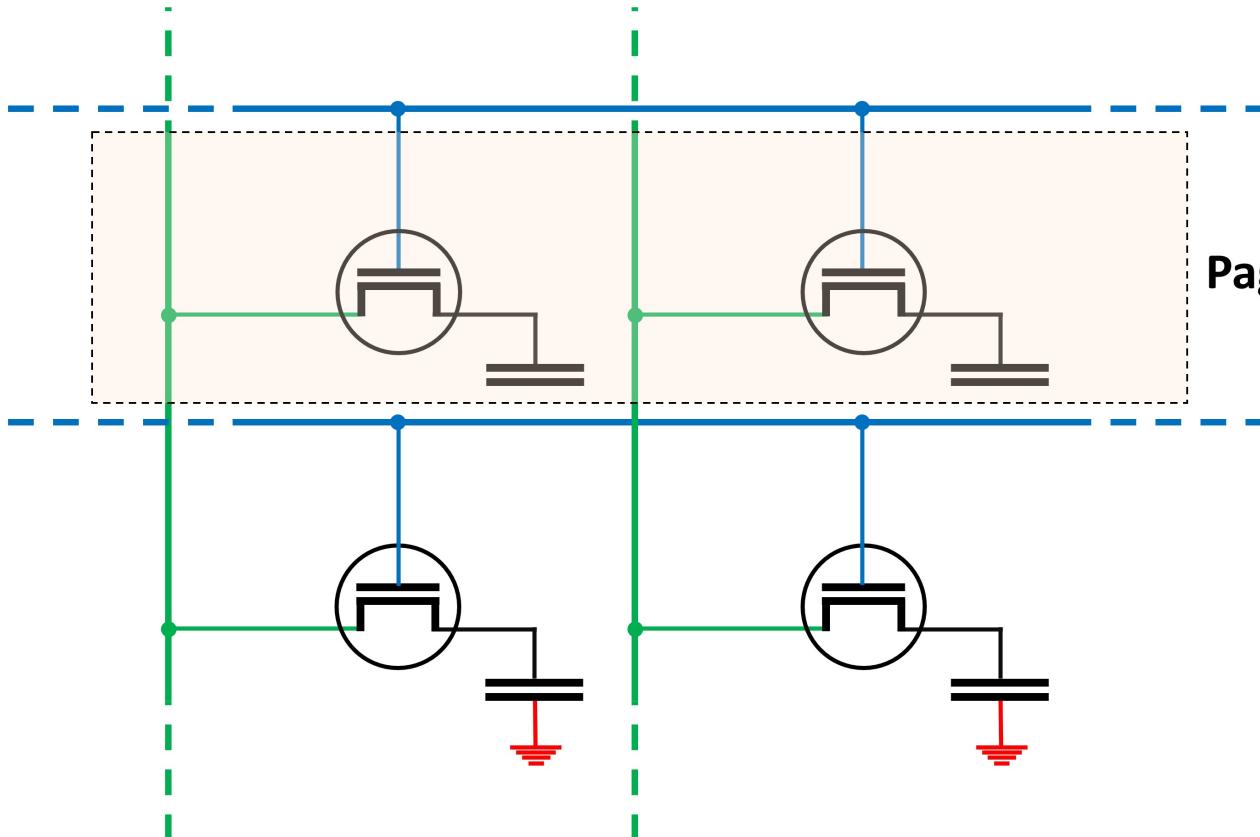


Data Cell

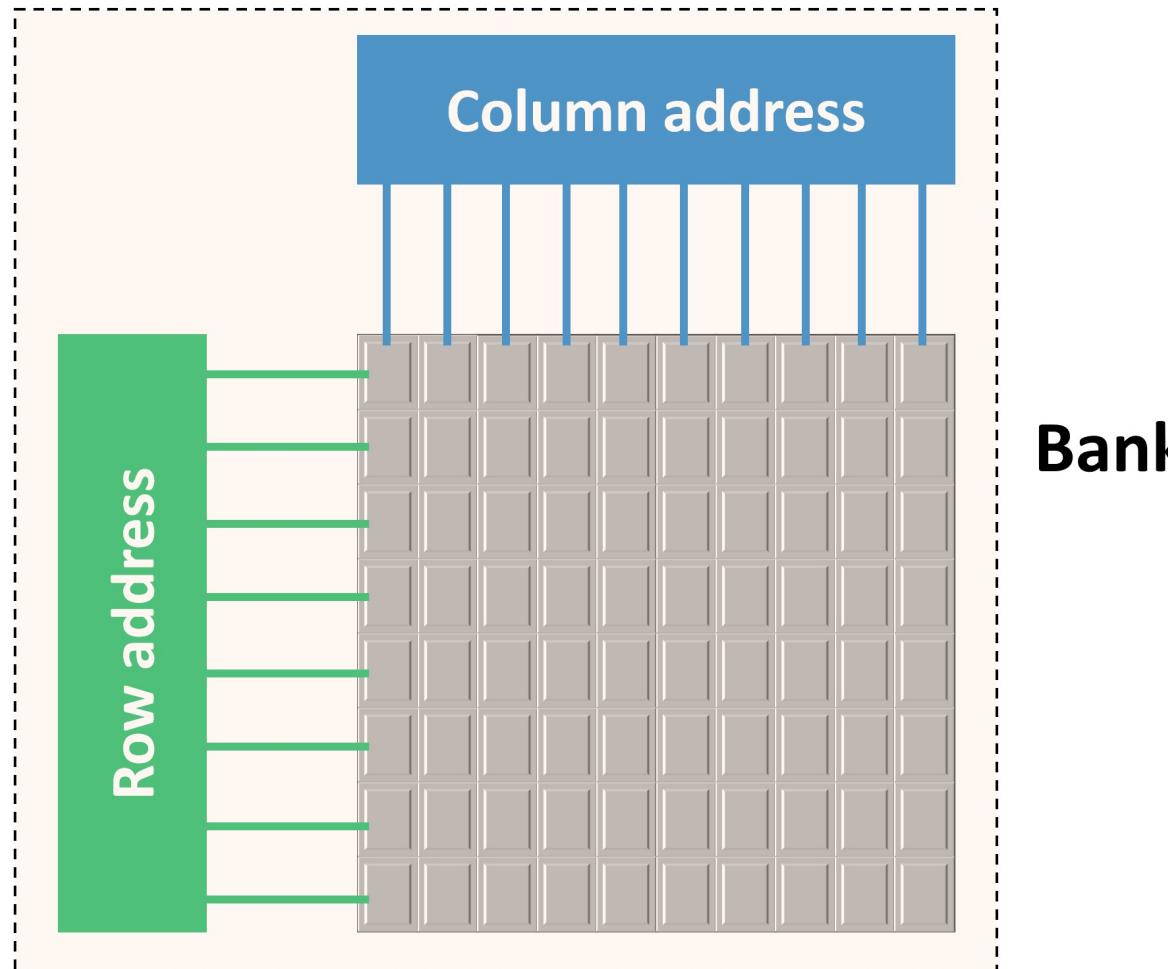


Page

Page

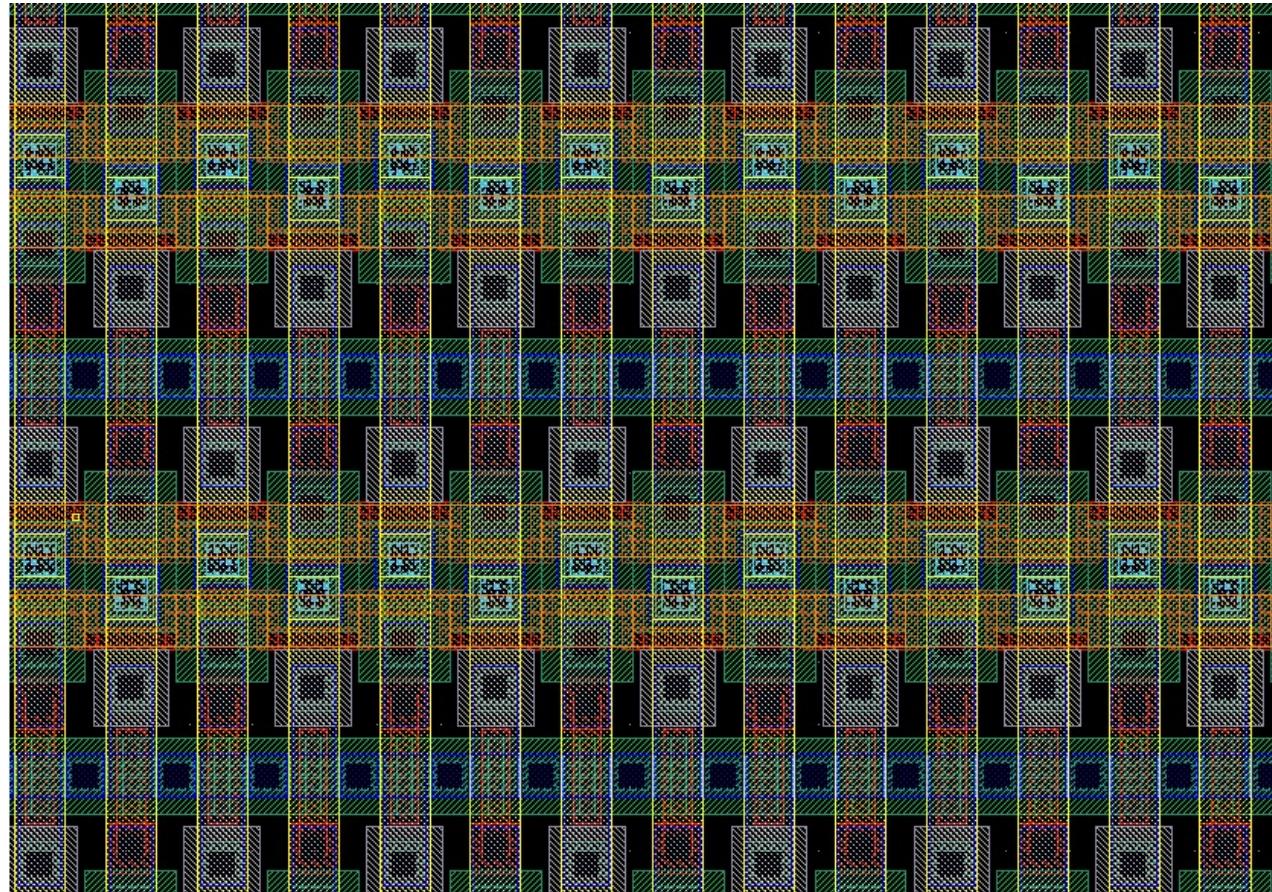


Bank

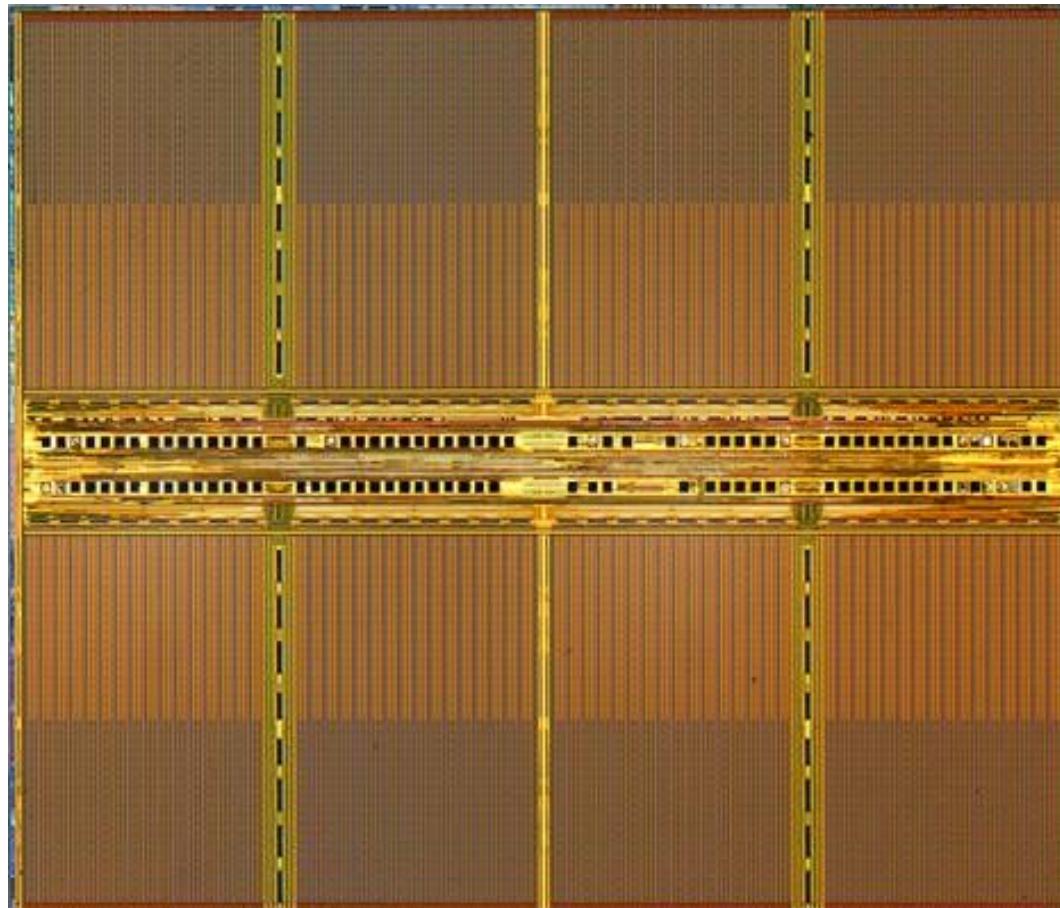


Bank

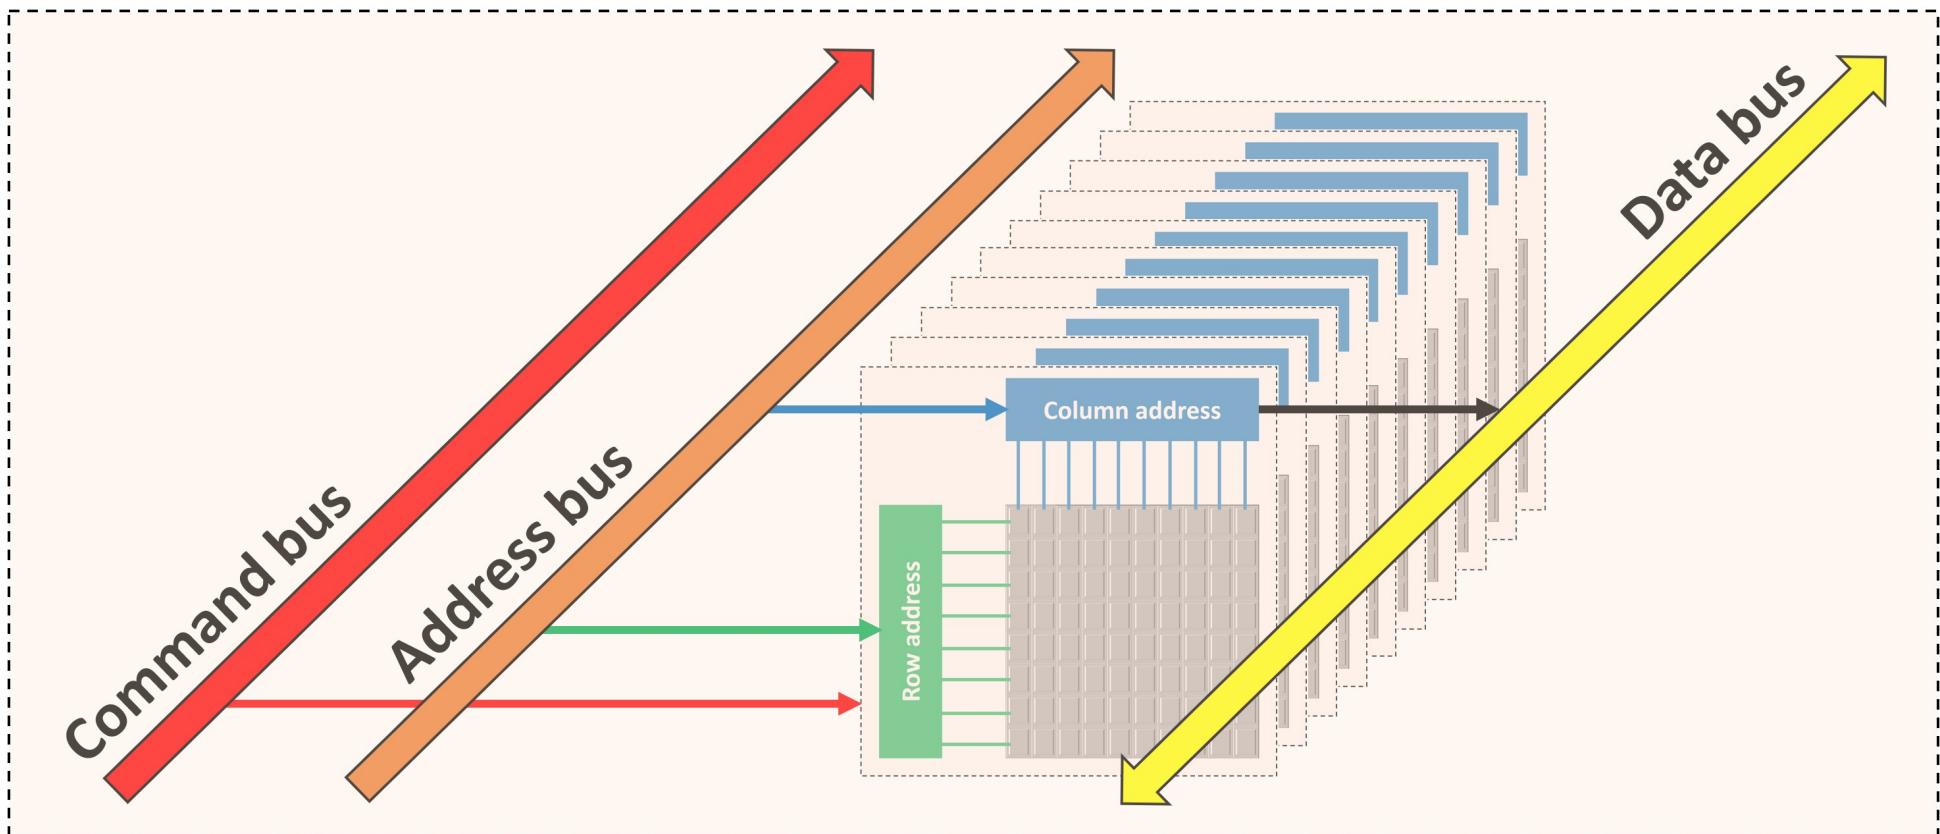
Layout of Memory Bank



Samsung's DRAM cell

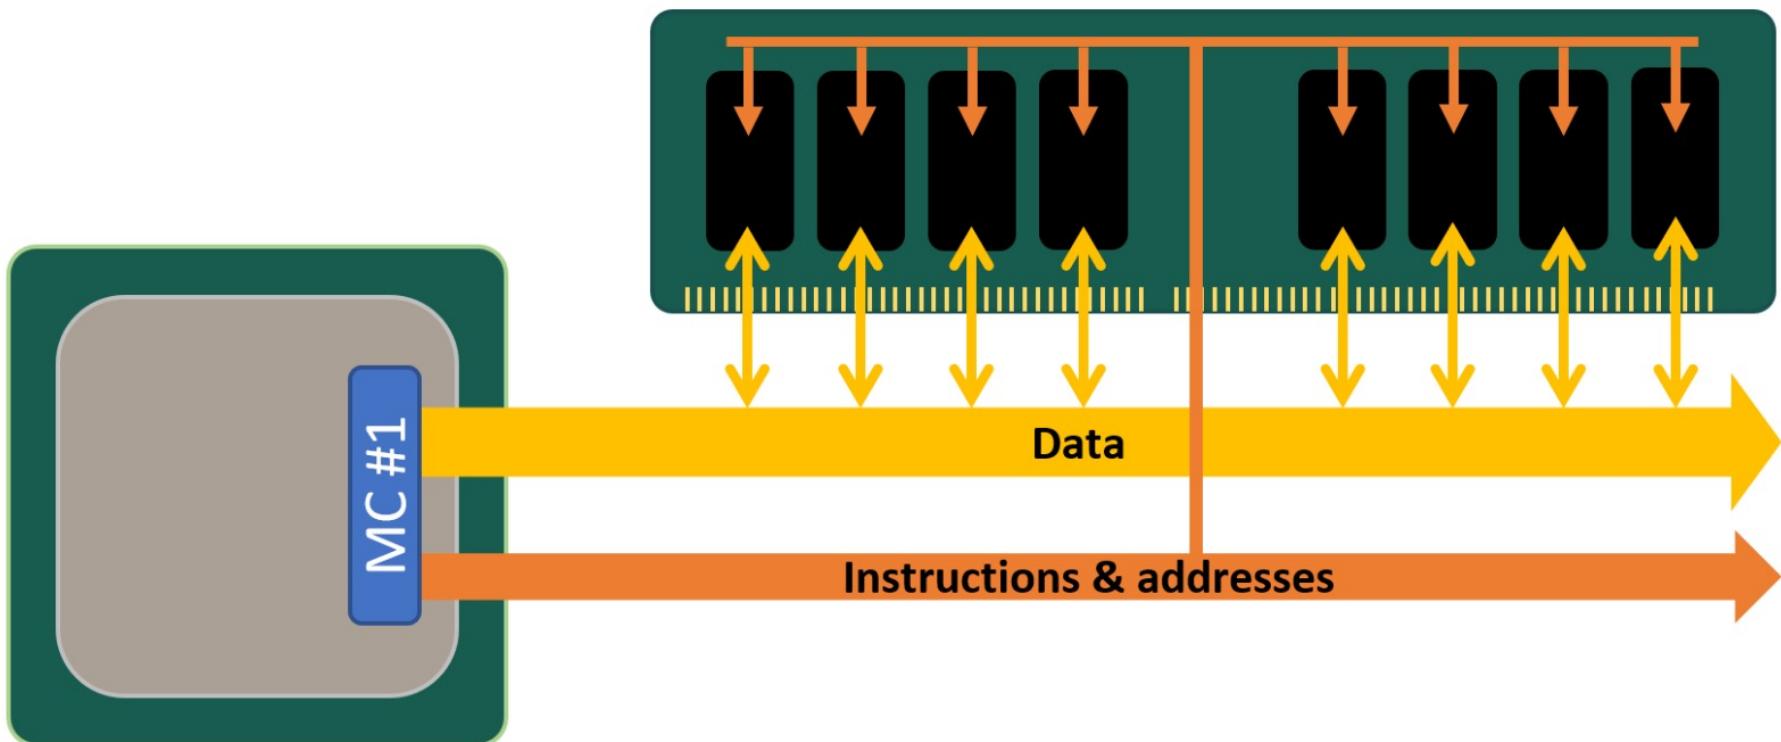


Module

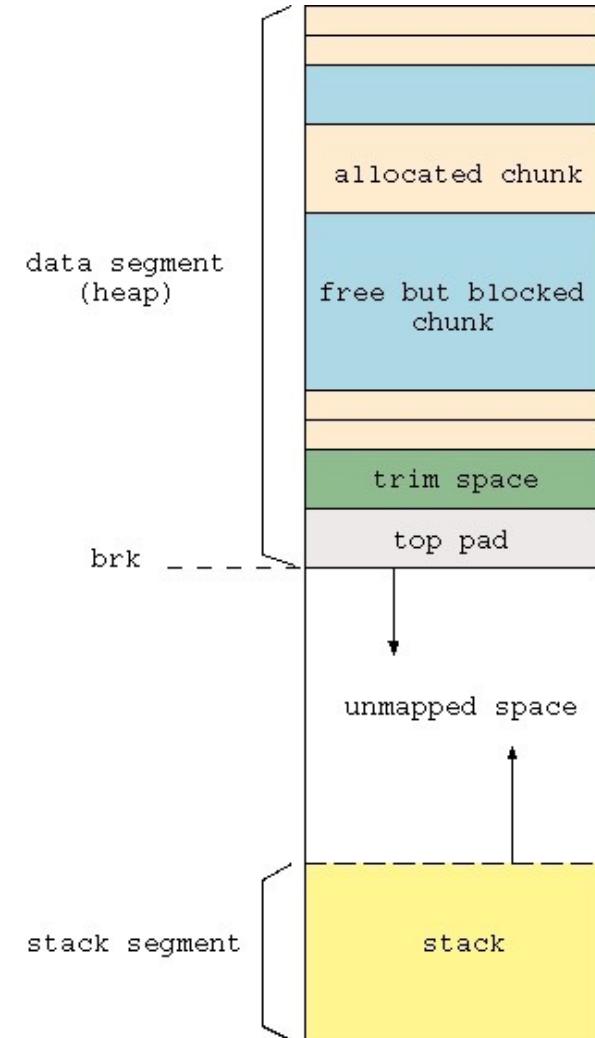
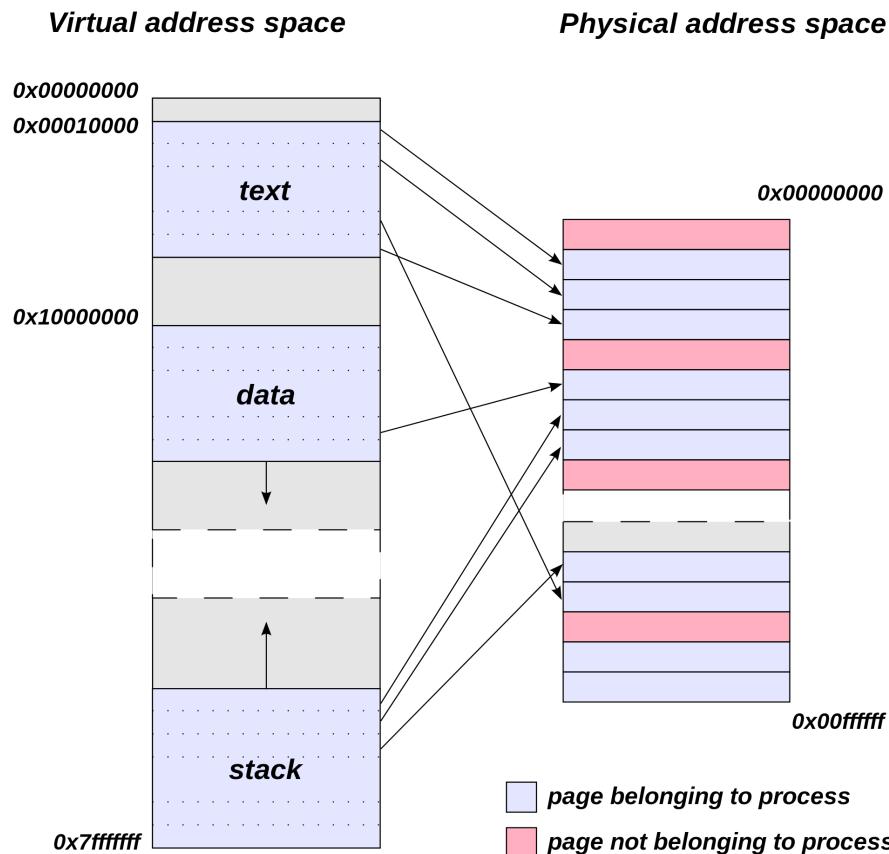


DRAM module



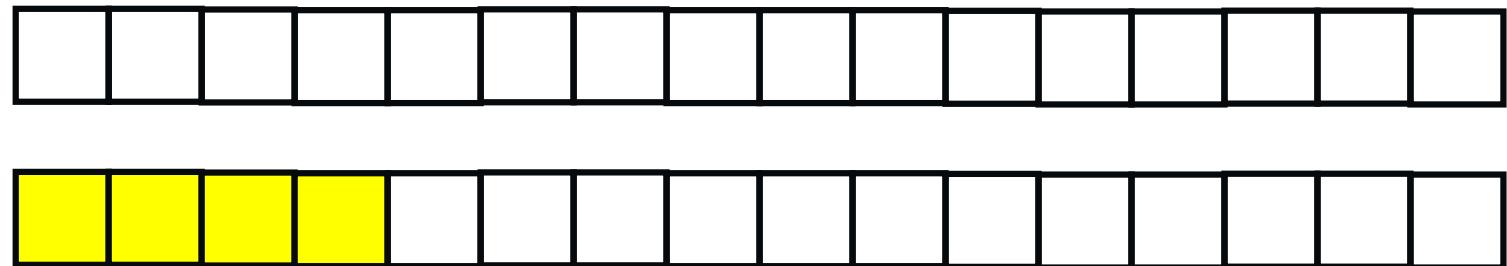


Virtual Address Space And Physical Address Space



Allocations

P1 = malloc(4)



P2 = malloc(5)



P3 = malloc(6)



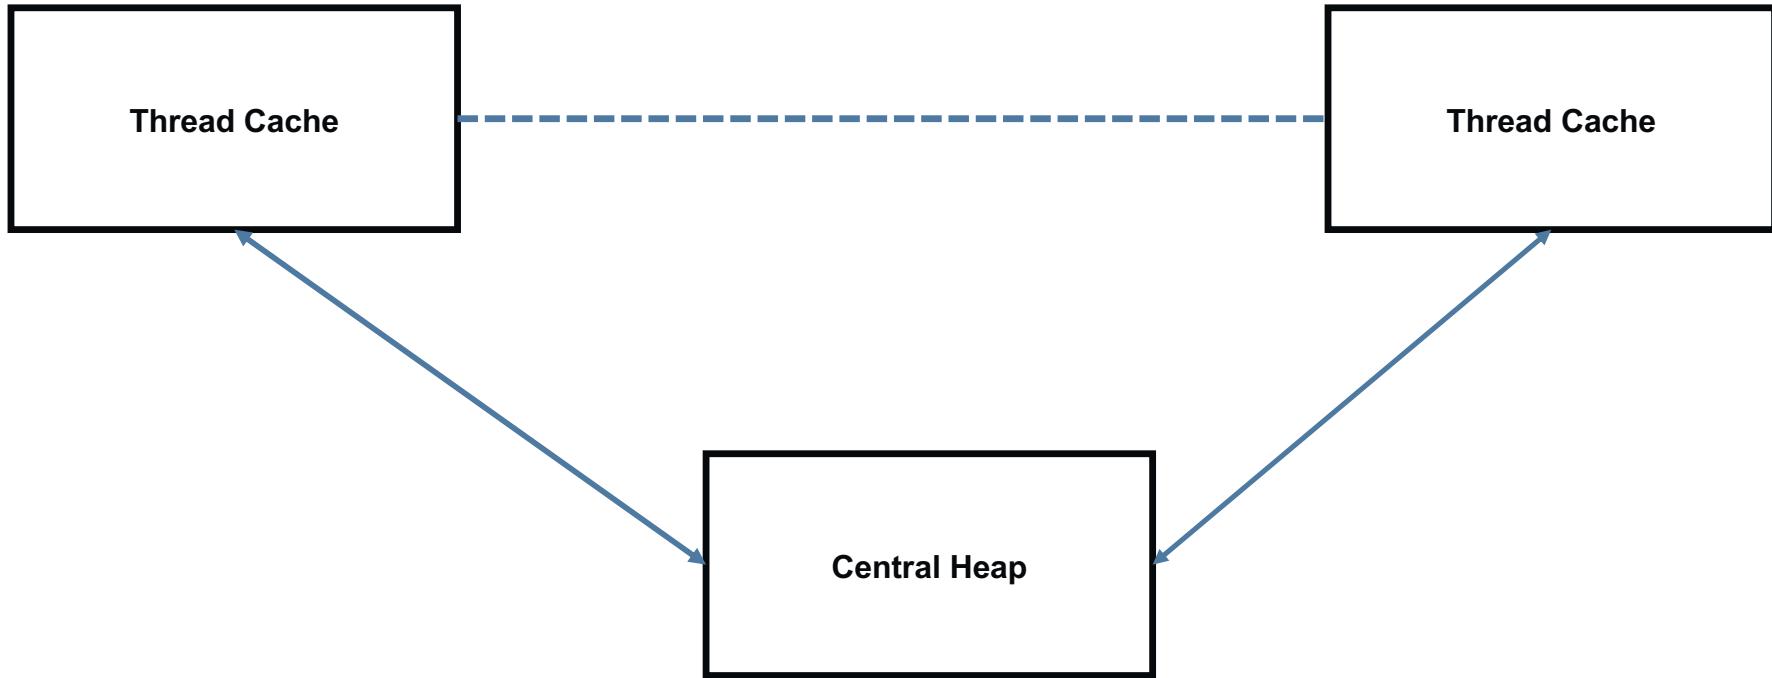
free(p2)

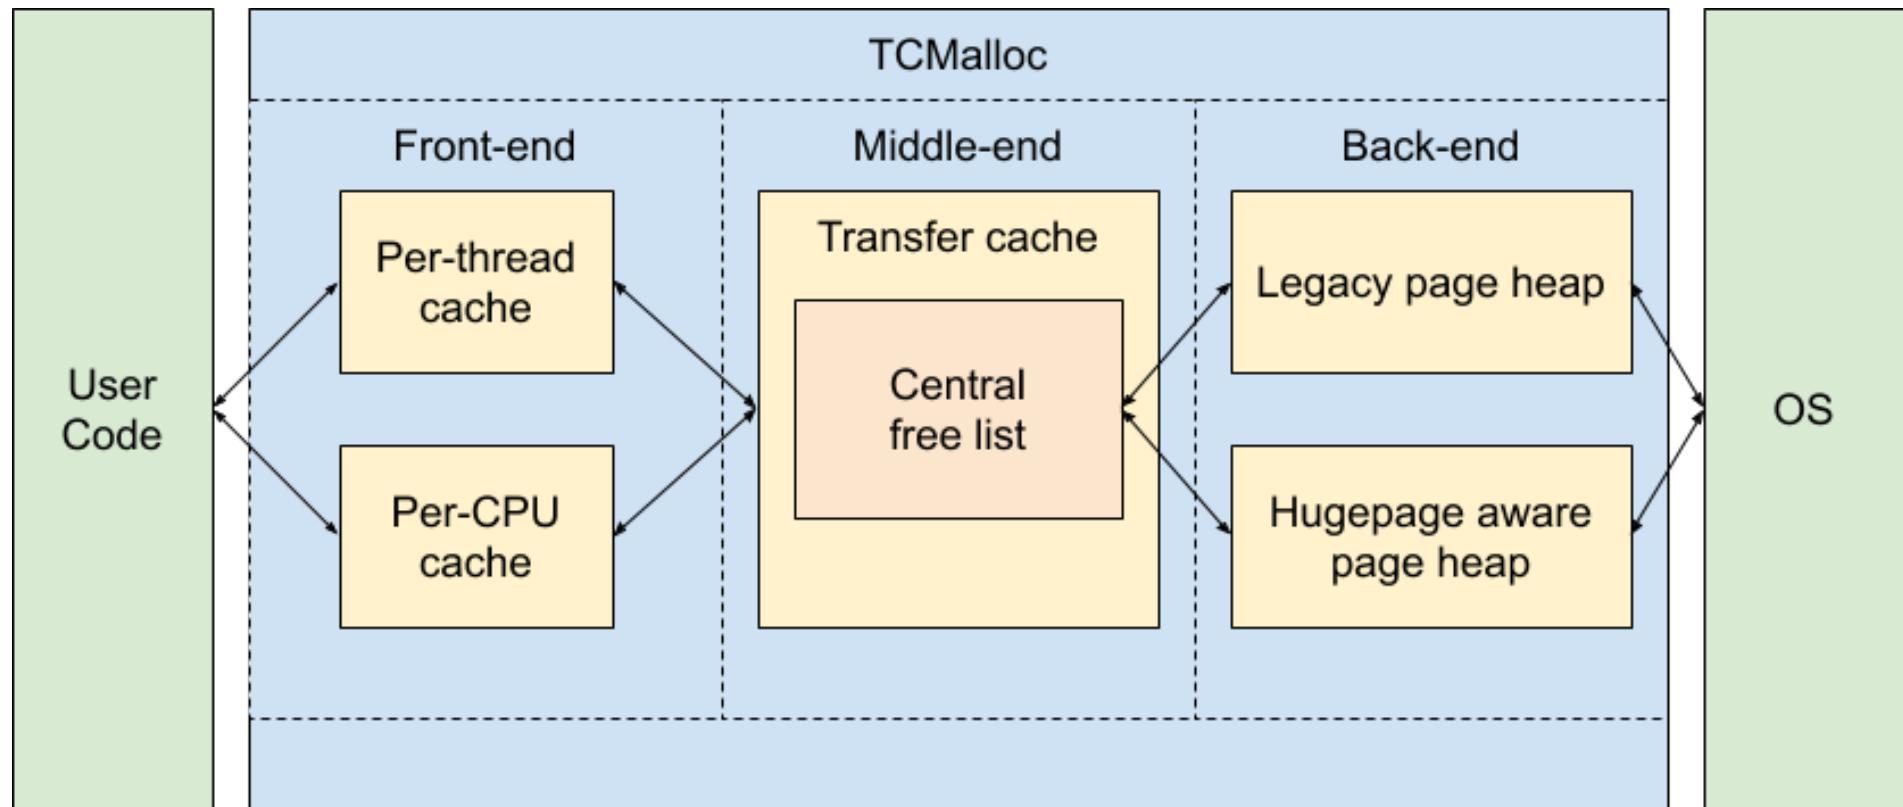


P3 = malloc(6)

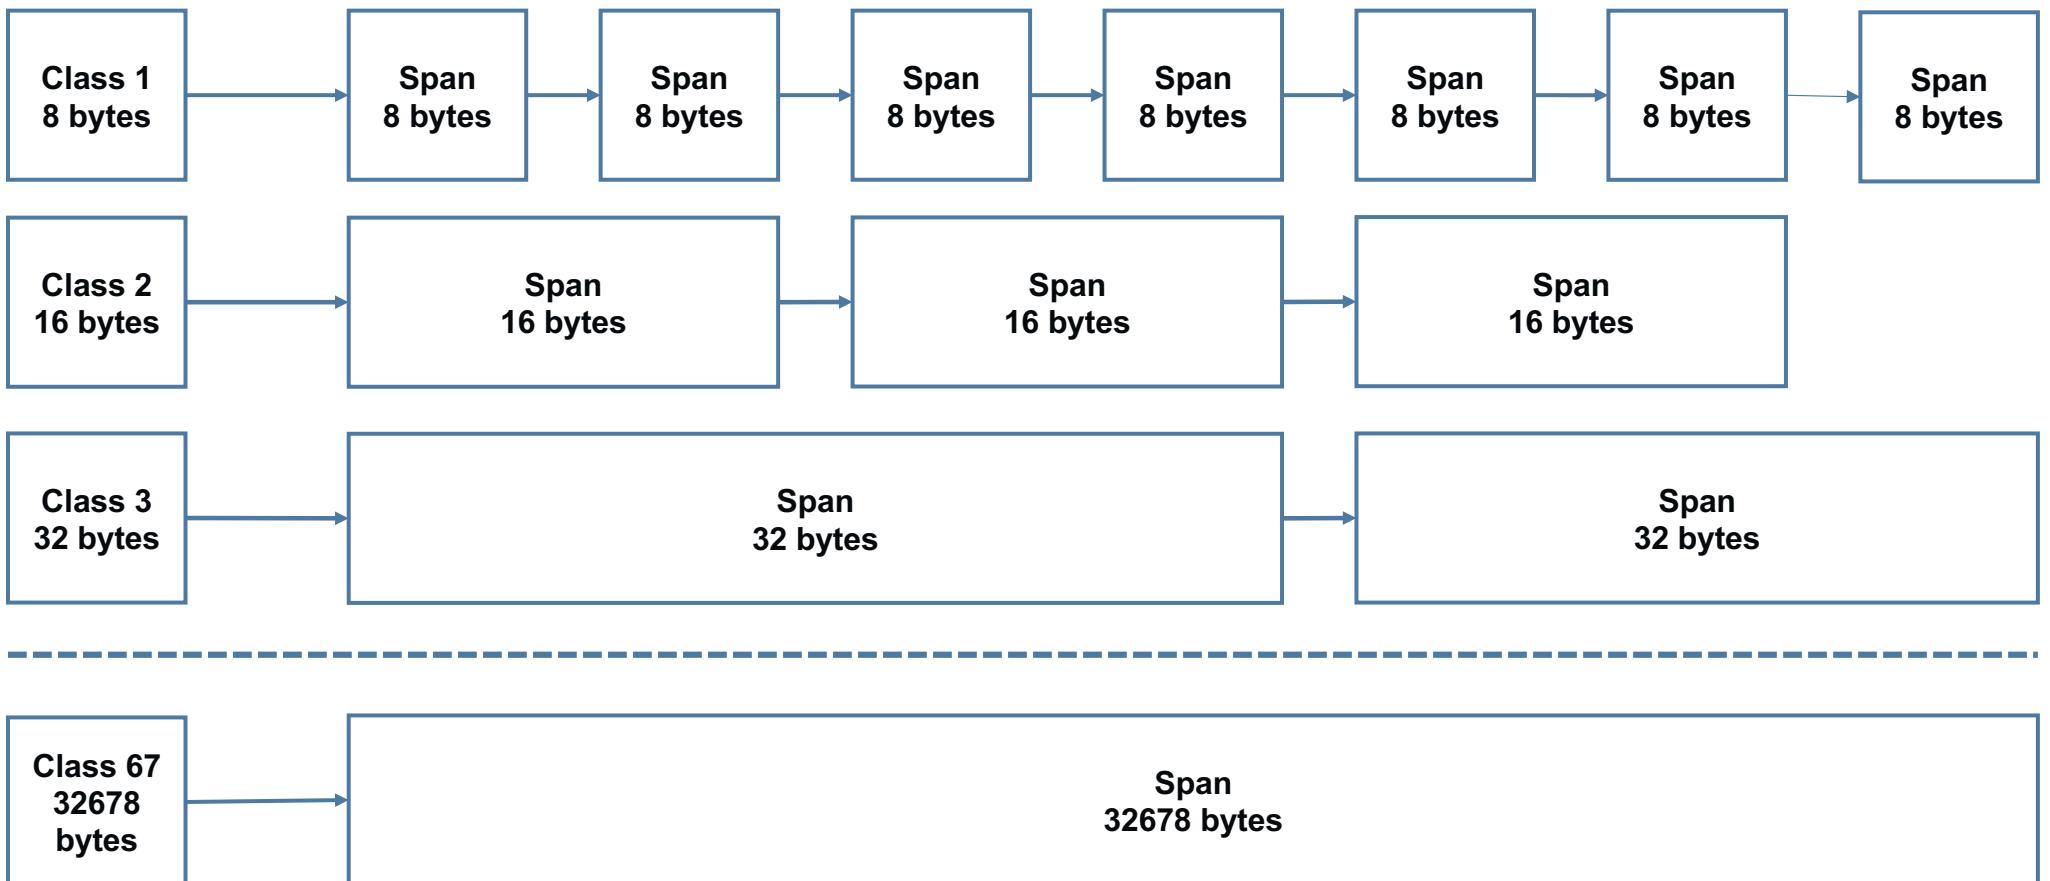
?

TCMalloc

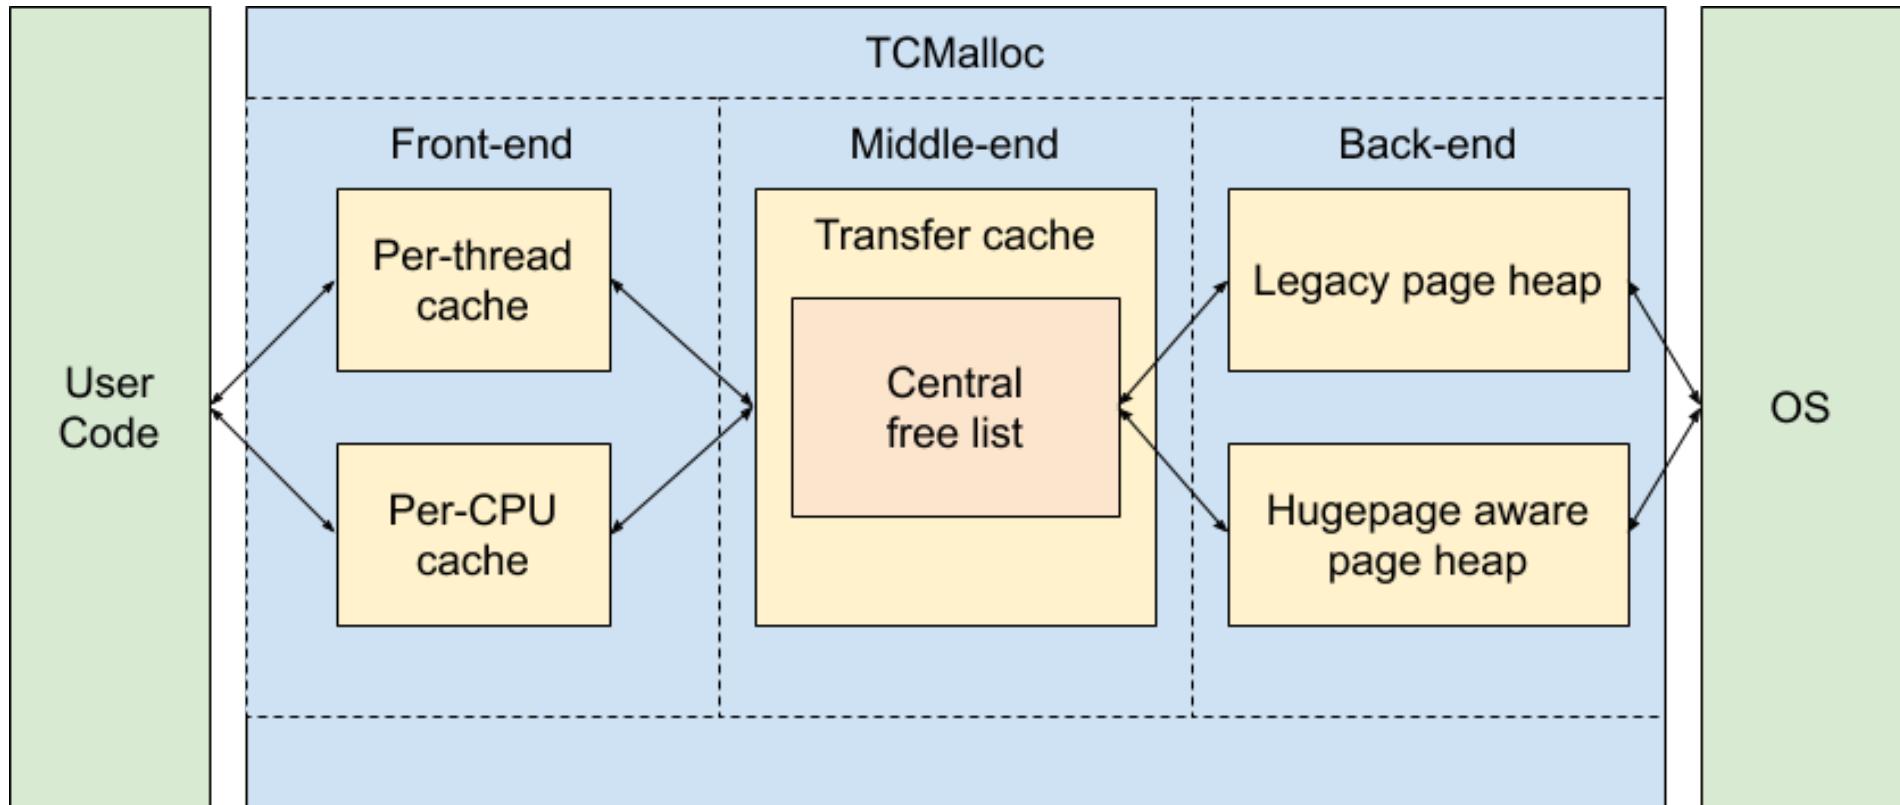




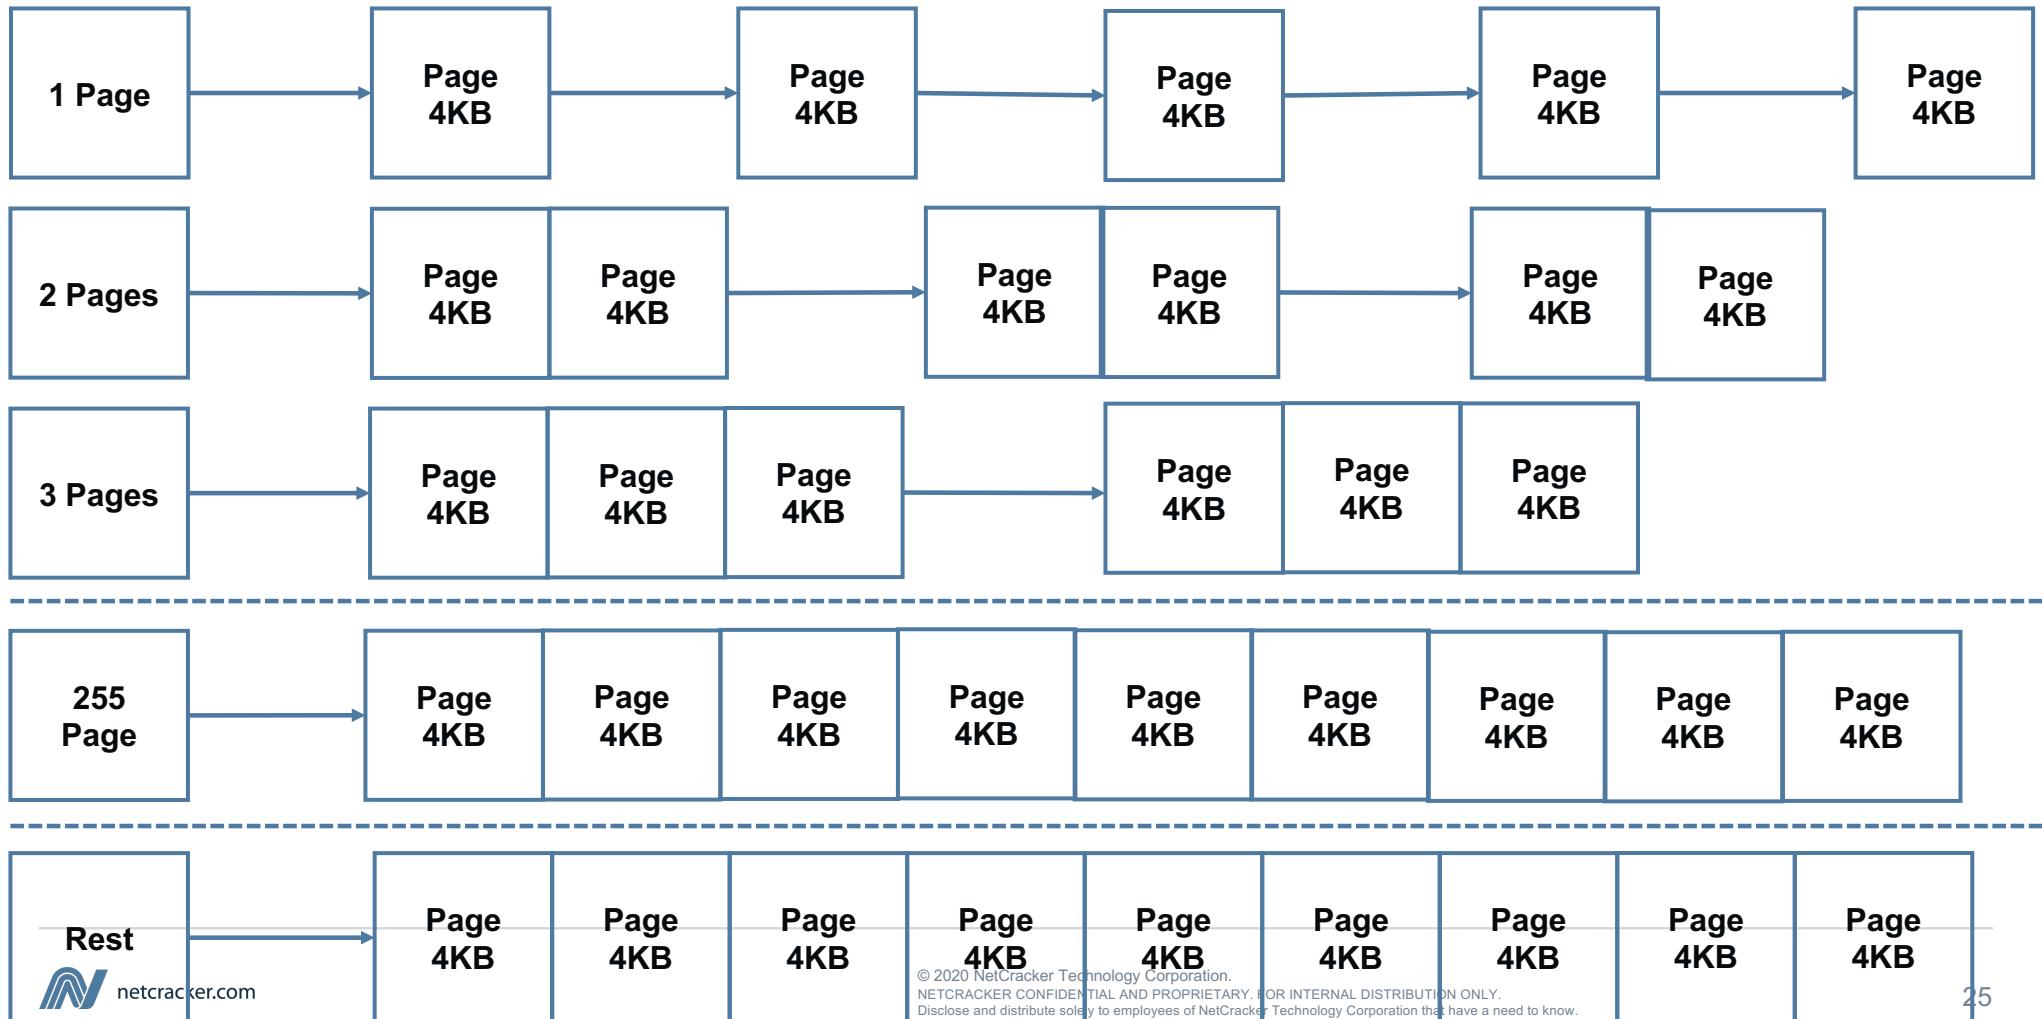
Thread Cache - Frontend



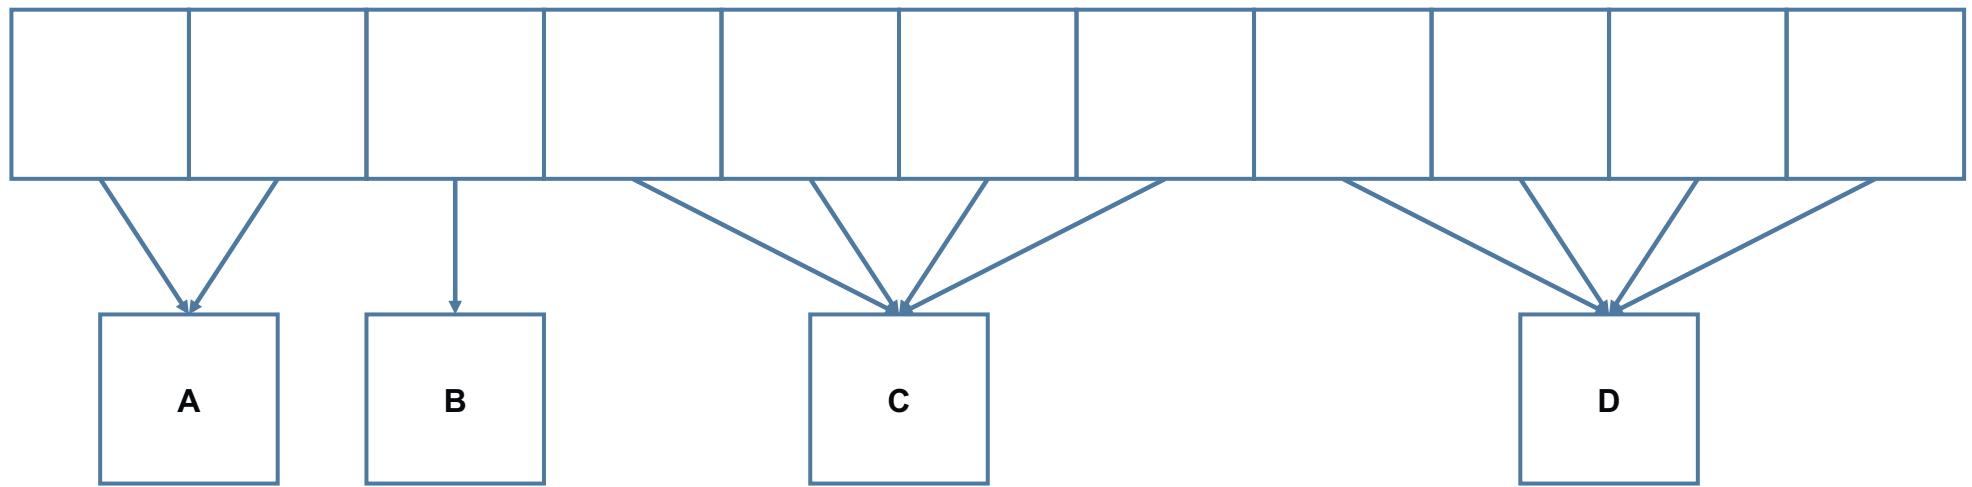
Middle End And Back End



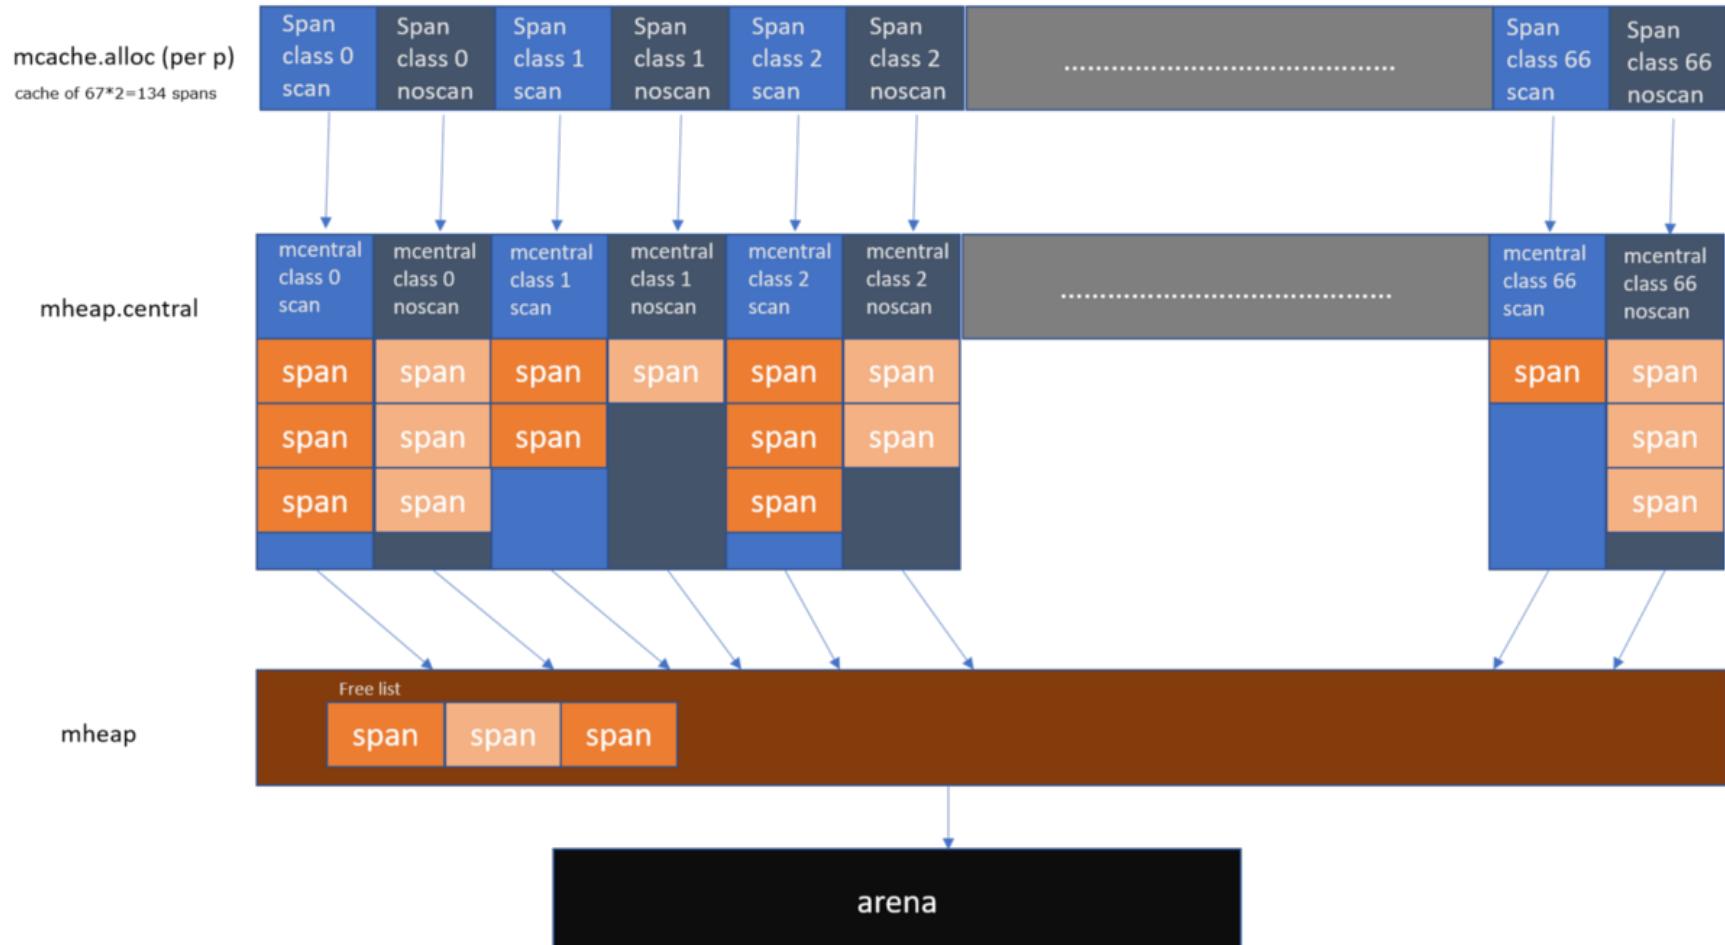
TCMalloc Backend



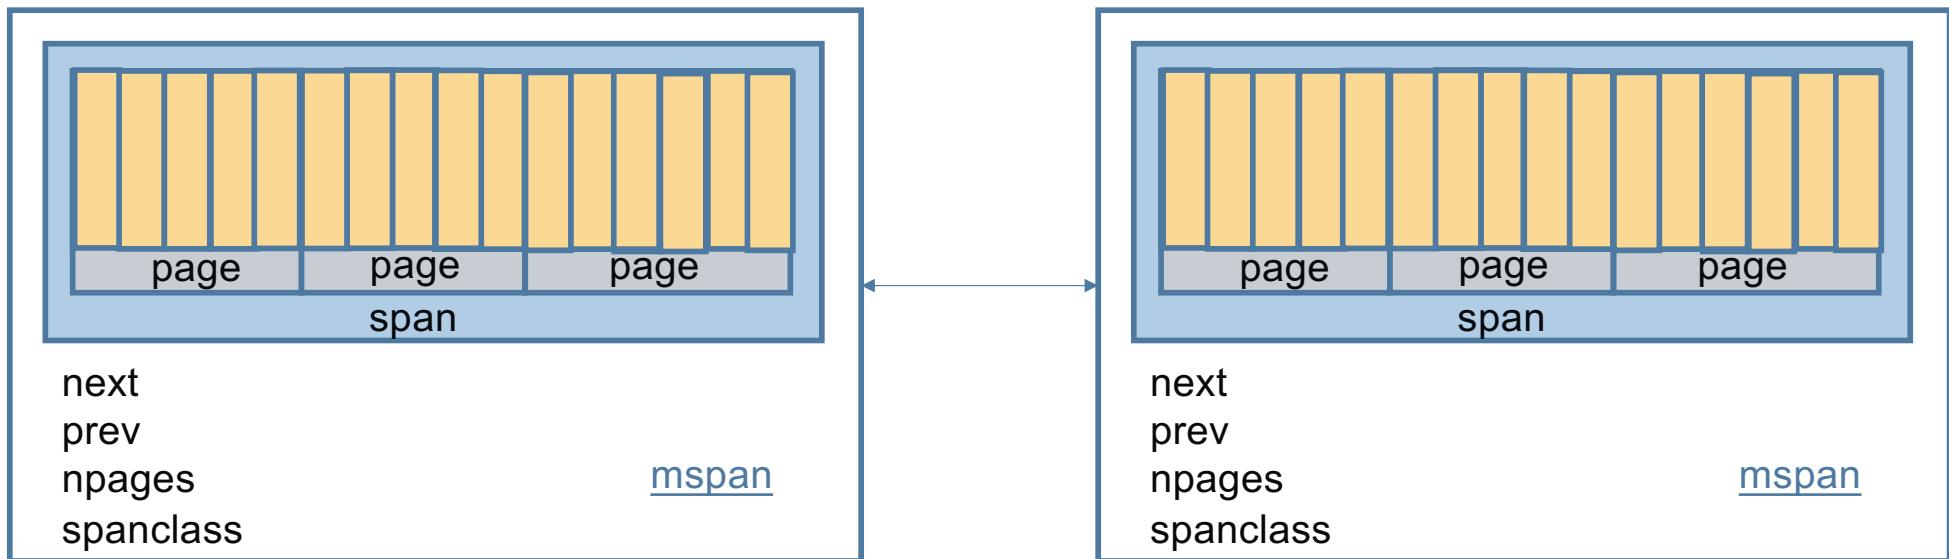
Spans



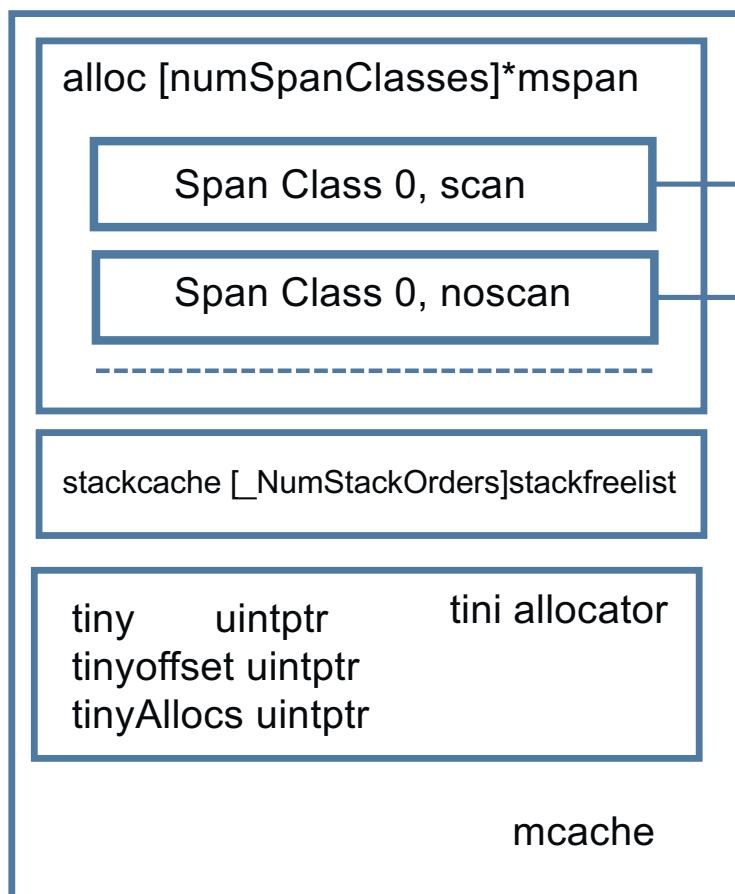
Memory Management

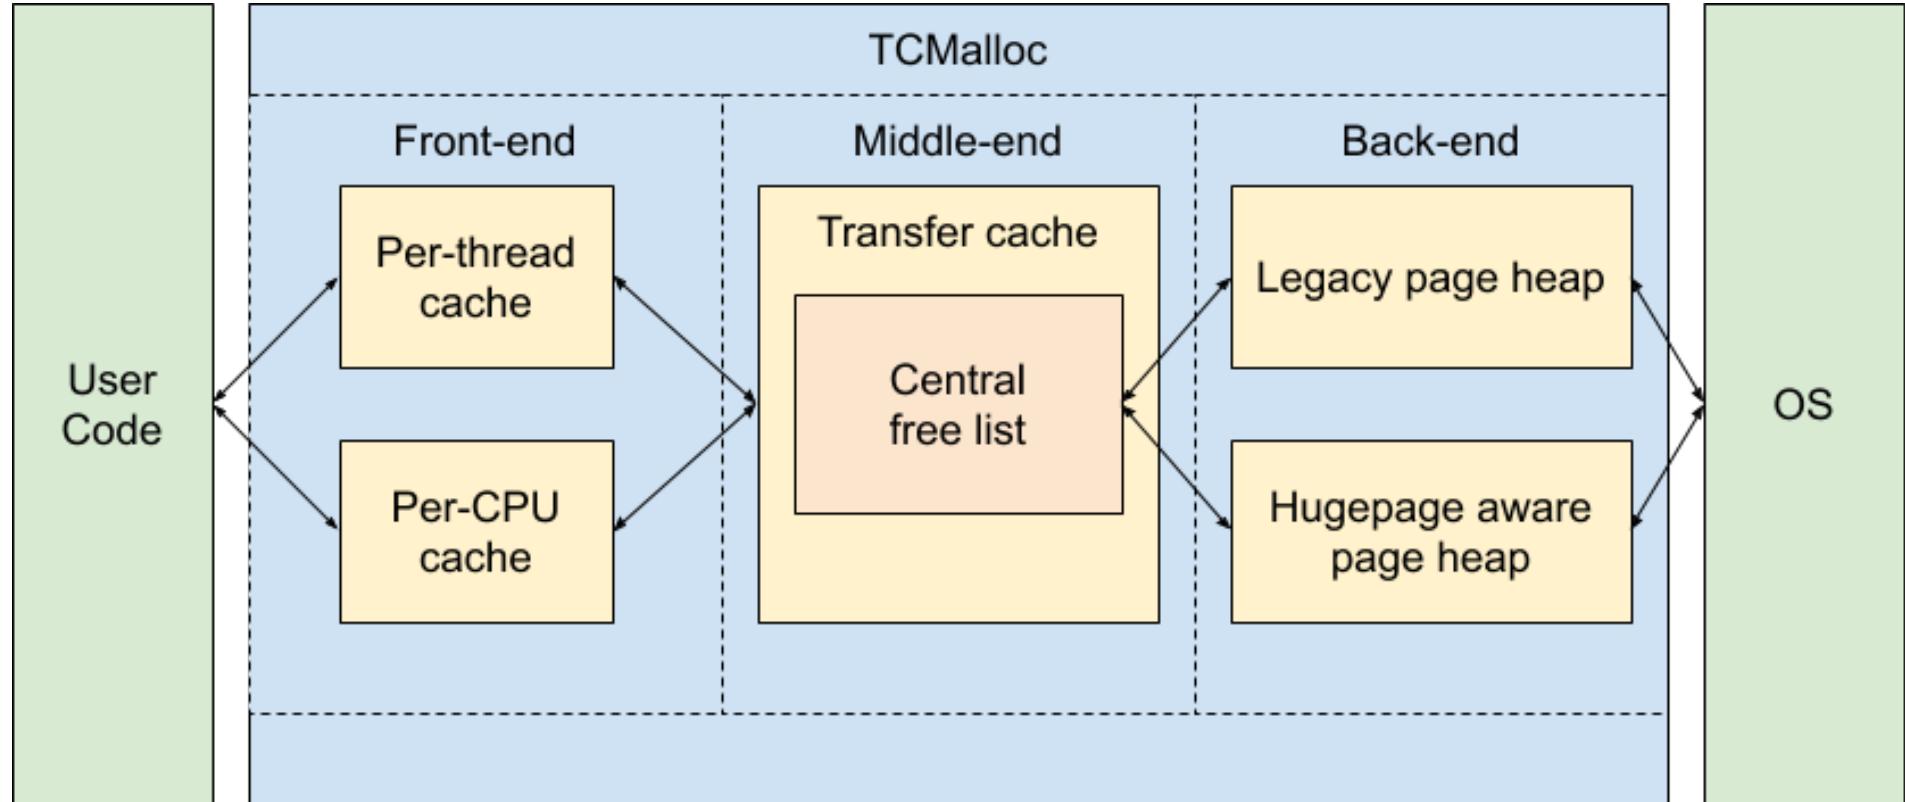


mspans

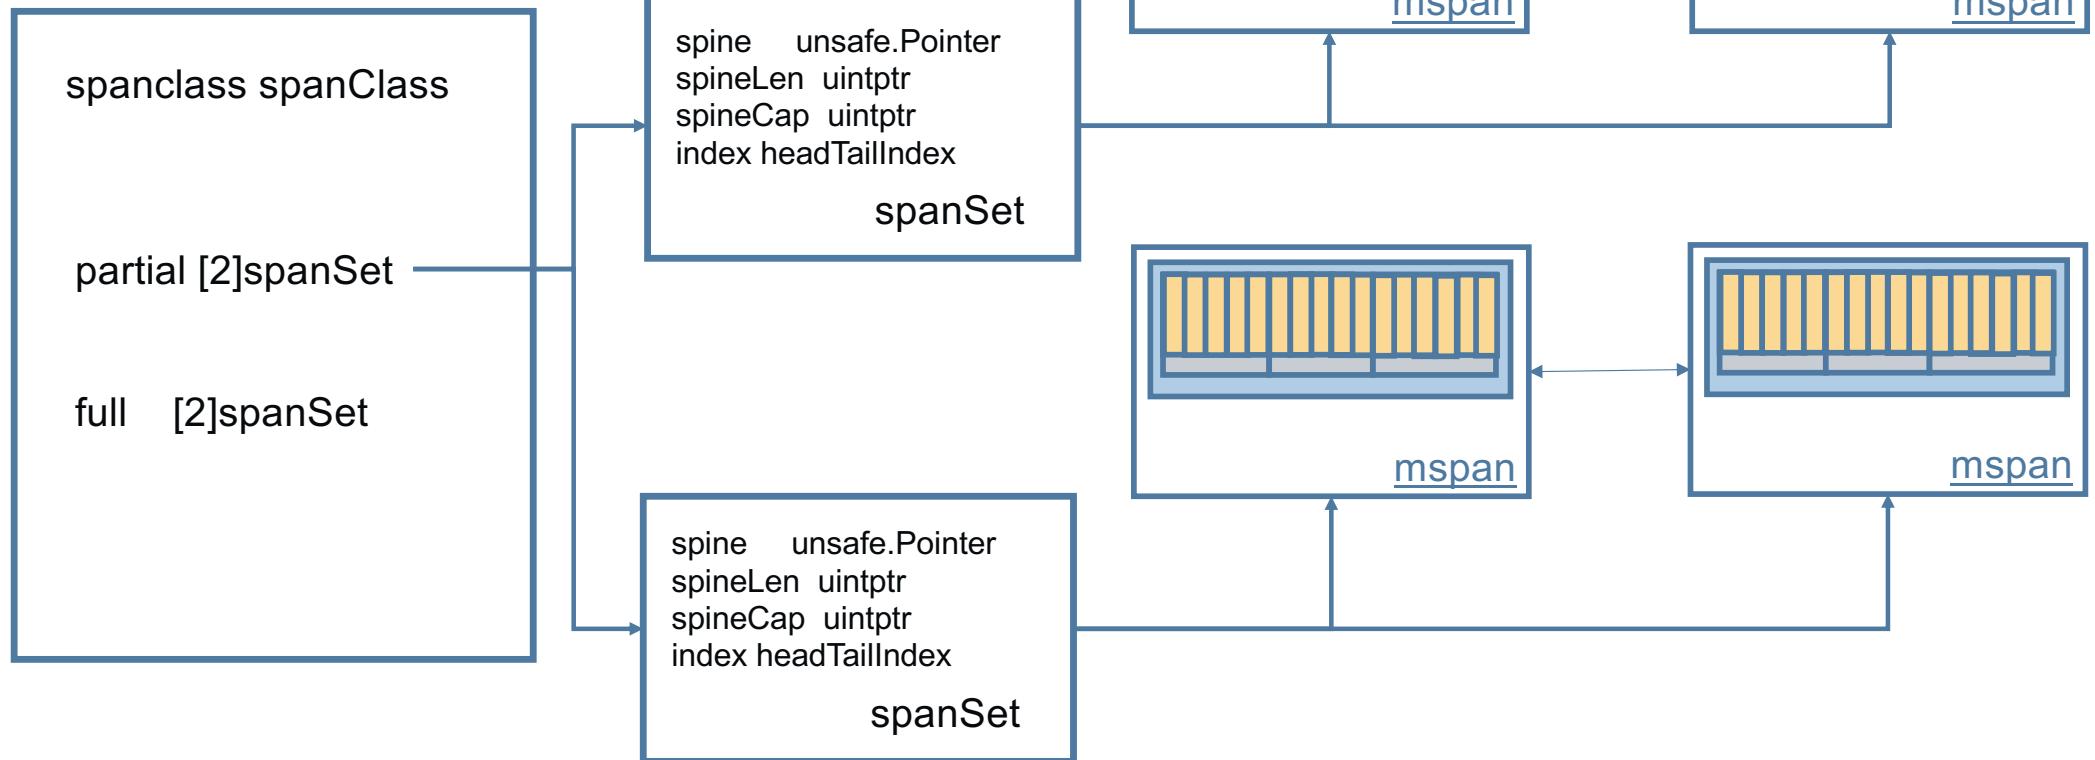


mcache

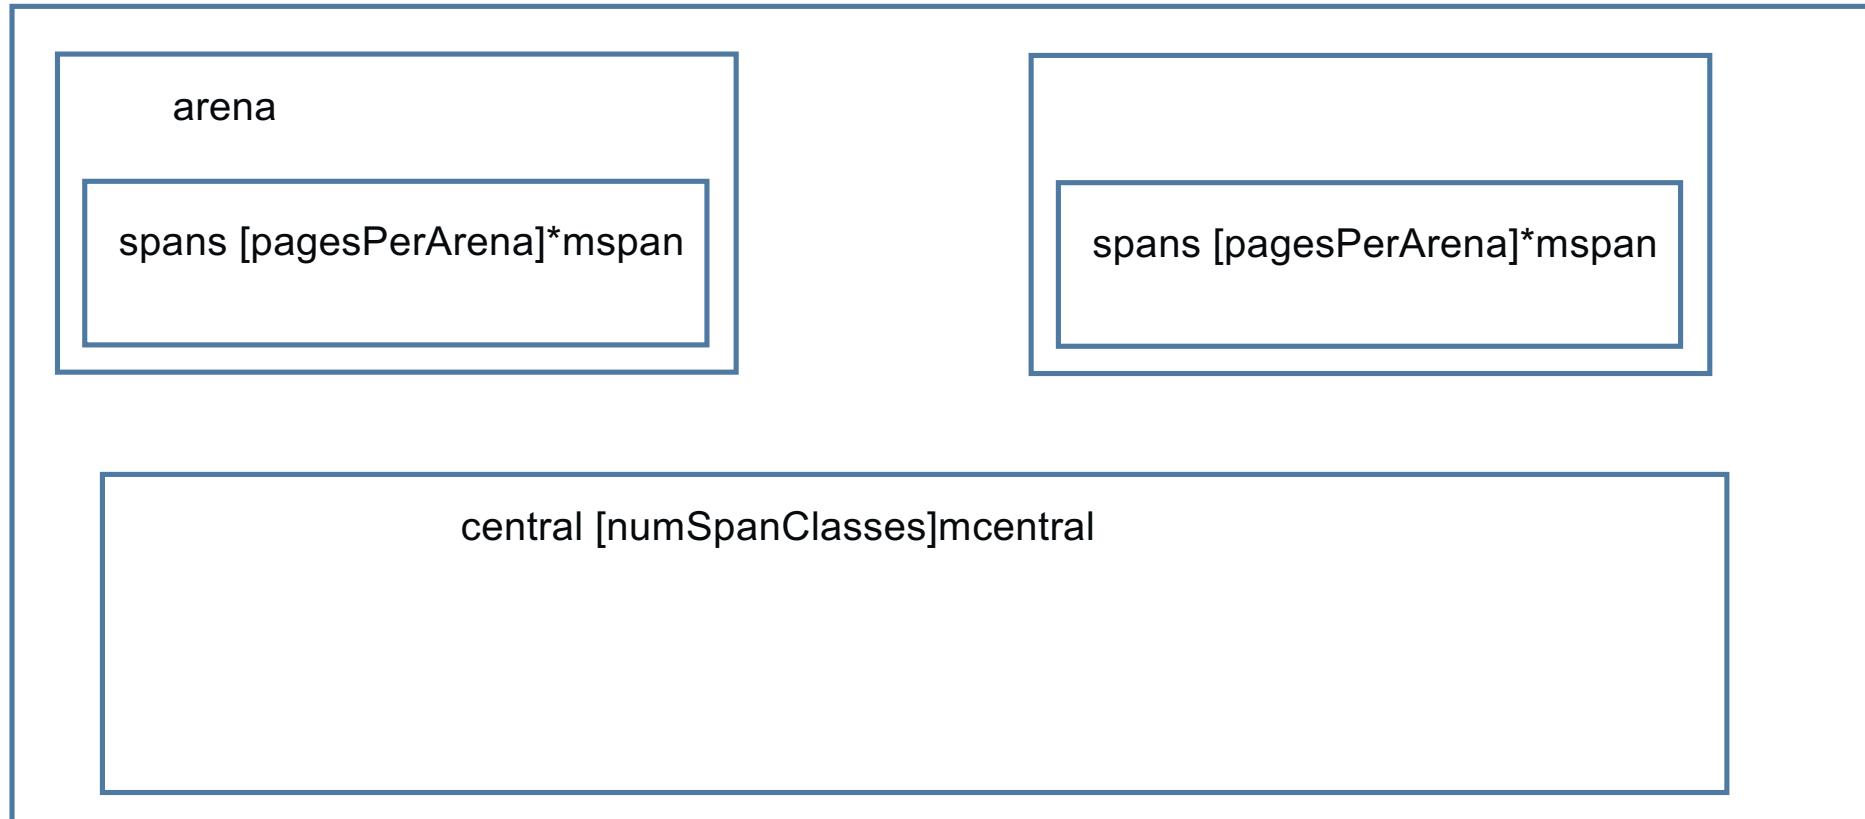


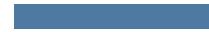


mcentral



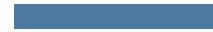
mheap





Memory allocation

- Large then 32 KB -> mheap
- Less then 16 KB -> mcache -> tiny allocator
- 16-32 KB -> calculate size class -> mcache -> mcentral -> arena -> OS



GC Types

Free At Exit, aka There's Plenty of RAM

- simplest possible system
- best concurrency (only allocator)
- reasonable for small programs
- definition of "small" increases as technology improves
- no destructors

Reference Counting

- synchronous - destructors useful
- accidentally amortized (but long pauses possible)
- simple concurrency
- possible to retrofit
- expensive in cpu
- needs extra word per object to hold counts
- no cycles
- complicated api and/or leaky abstraction
- expensive allocator (fragmentation, locality)
- threading problems (mutating ref counts - no read-only data!, destructor runs on random thread)
- iOS, file systems

Mark-Sweep Collector

- traversal required
- simple
- destructors easy, but delayed
- conservative option (traversal can be approximated)
- possible to retrofit
- asynch
- expensive allocator (fragmentation, locality)
- complicated concurrency (multi-color)
- Lua, Flash, Ruby, Go
- extensions:
- deferred free
- mark tables (examine multiple objects at once)

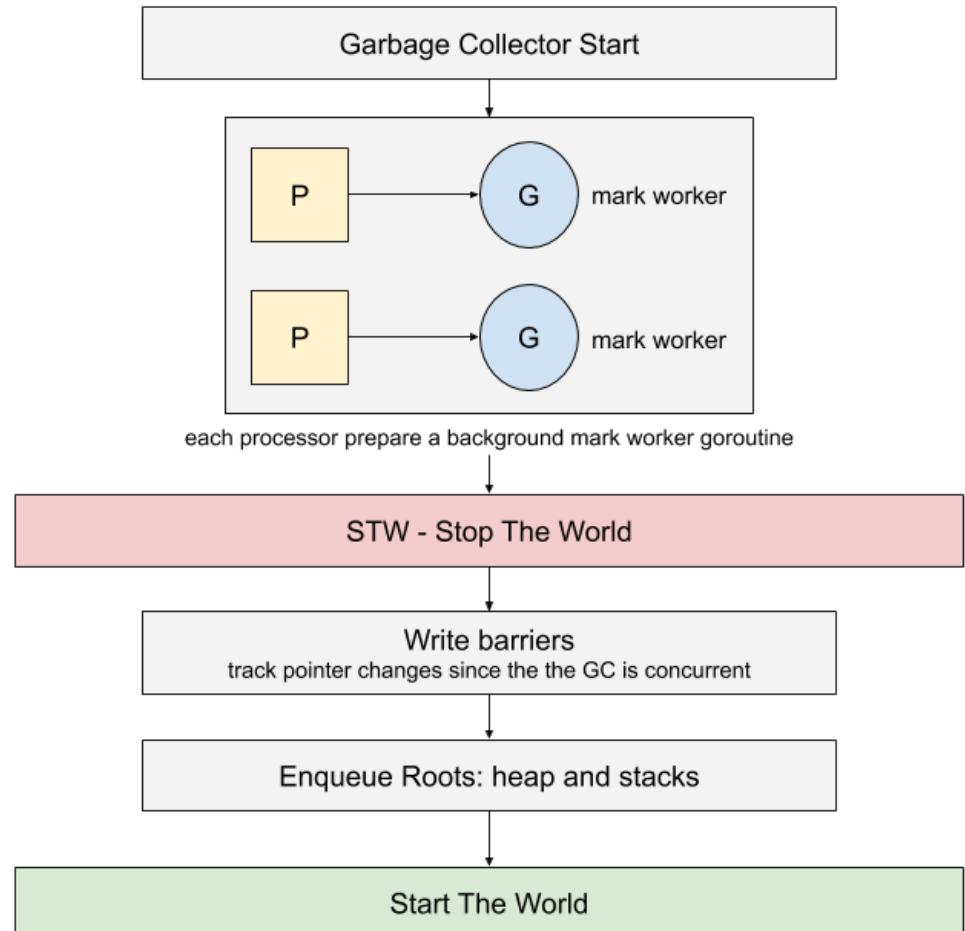
Mark-Compact Collector

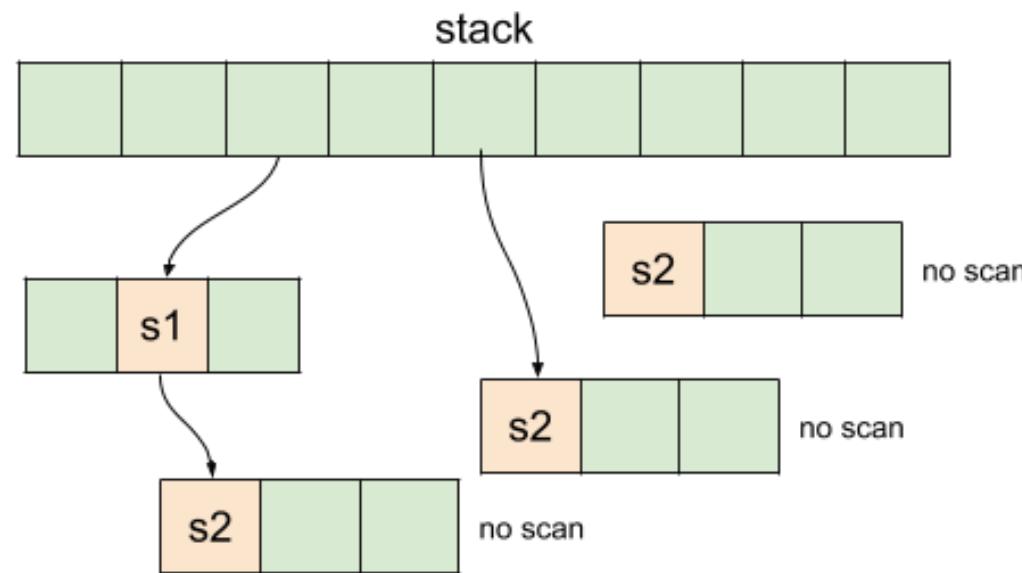
- precise traversal required
- simple
- minimum total memory usage
- super cheap allocator ("bump" allocator with only a few instructions)
- moving objects difficult to retrofit
- needs 3 passes, but can trade memory for performance
- very complicated concurrency (multi-color, barriers)
- general algorithm that can make other problems easier
- example: fair random row selection

Copying Collector

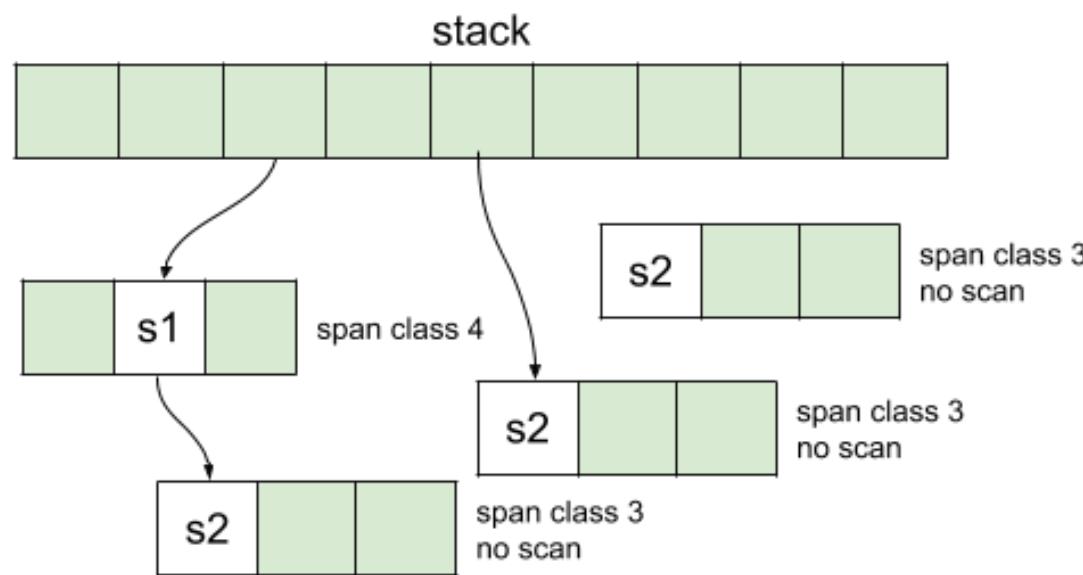
- precise traversal required
- simplest
- super cheap allocator
- work proportional to live data (garbage doesn't matter!)
- semi-spaces
- no destructors - finalize should not be used
- very complicated concurrency (multi-color, barriers)
- moving objects difficult to retrofit
- common degenerate case: per-transaction pool, delete when done
- most useful gc algorithm whenever you have little live data
- non-memory example: web session storage deletion on Amazon SimpleDB

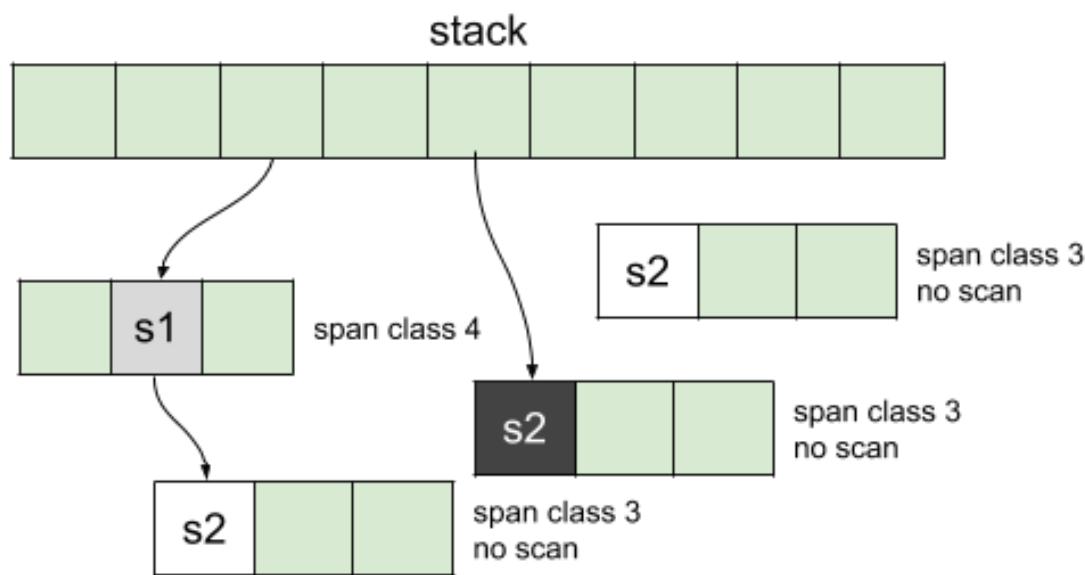
Collector Behavior

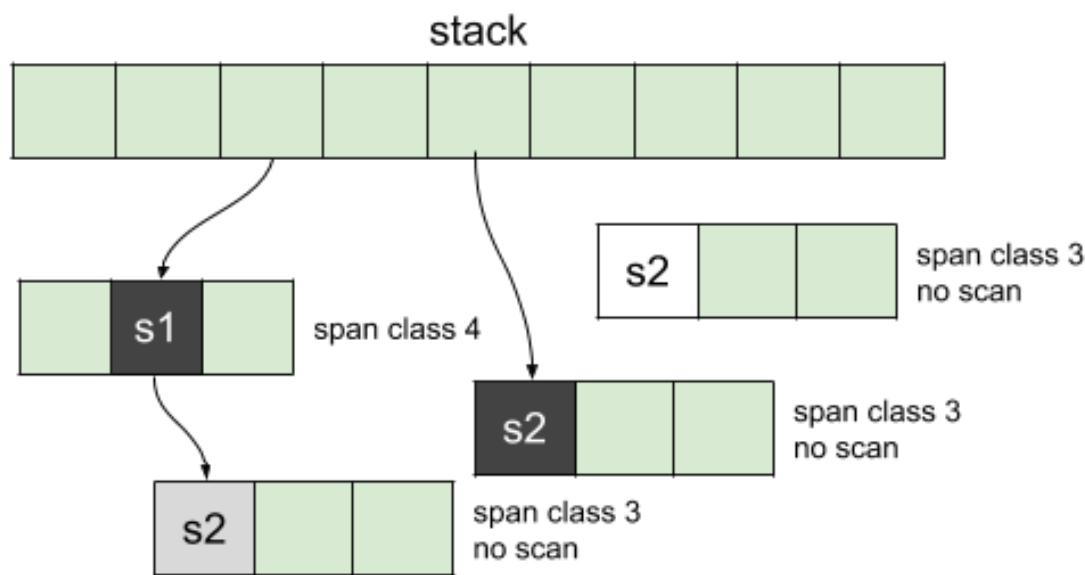


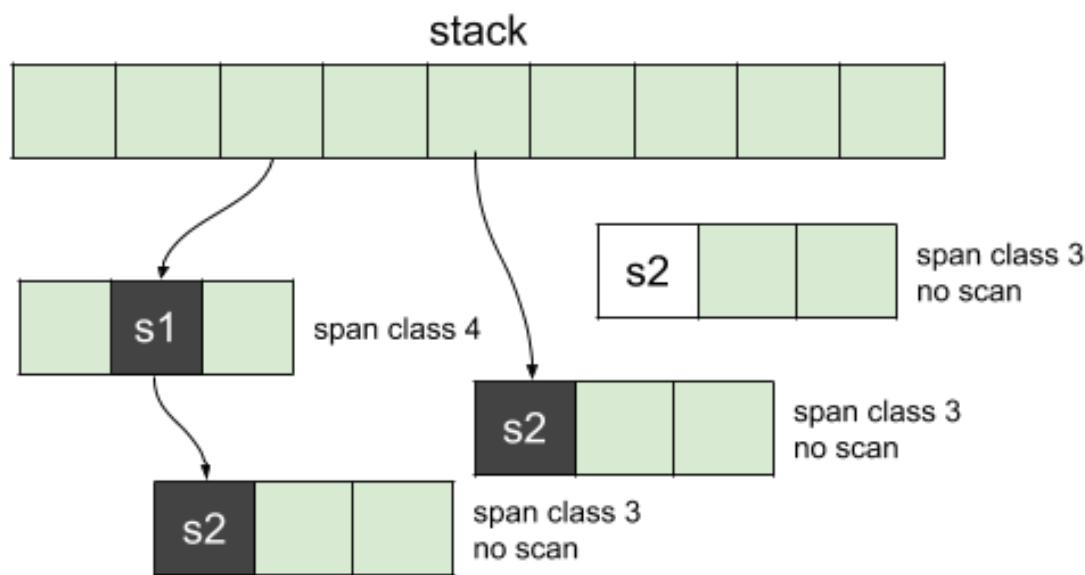


Coloring

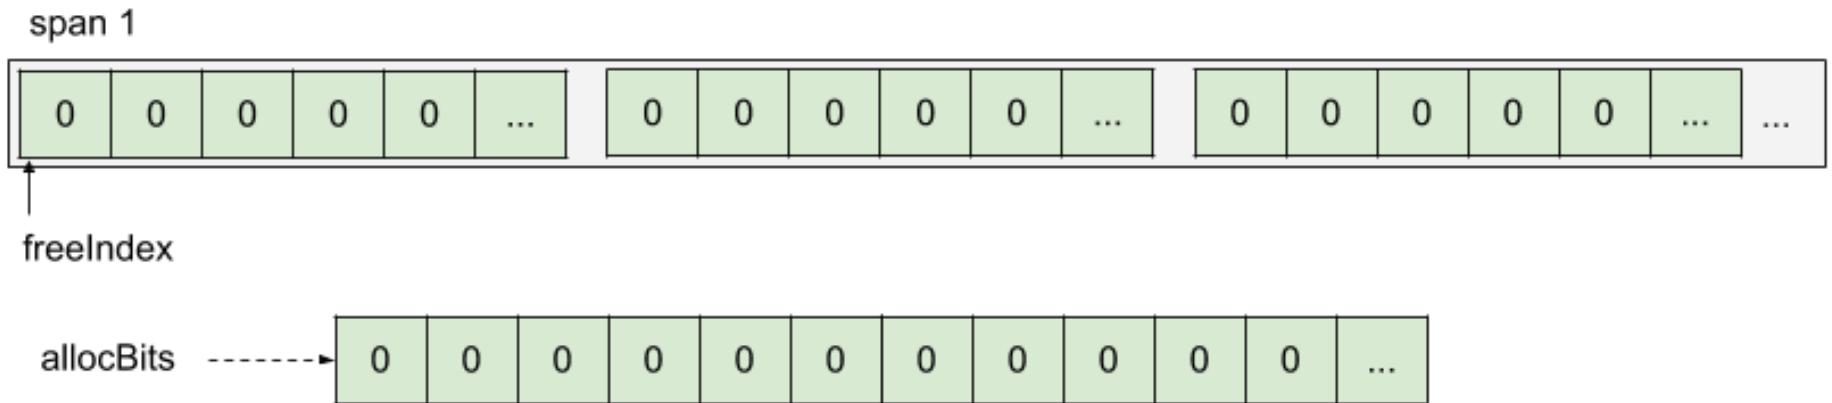


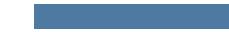




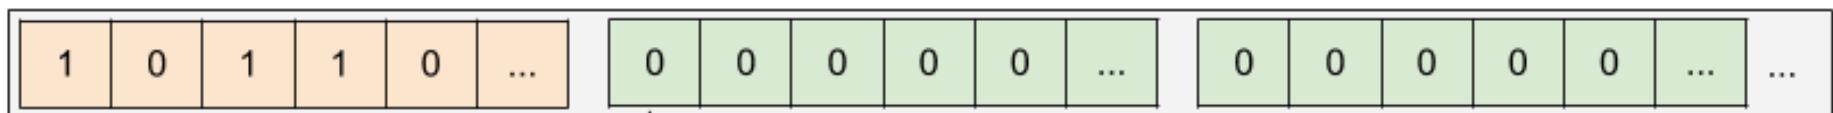


Bitmaps

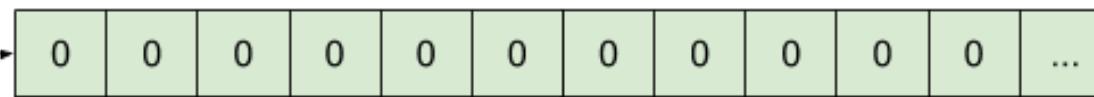


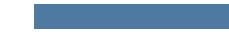


span 1



allocBits



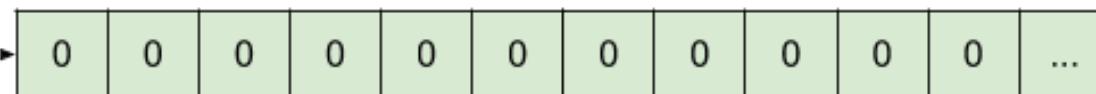


span 1

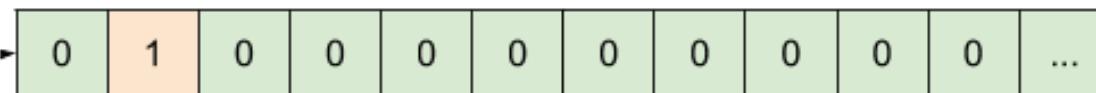


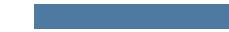
freeIndex

allocBits

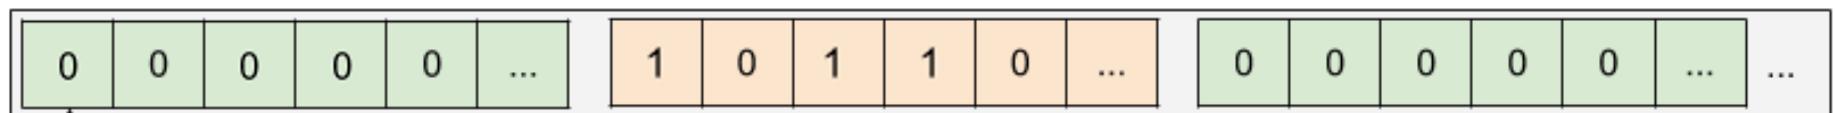


gcmarkBits



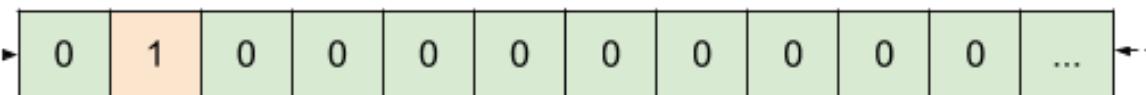


span 1

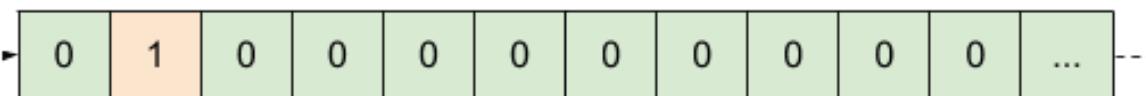


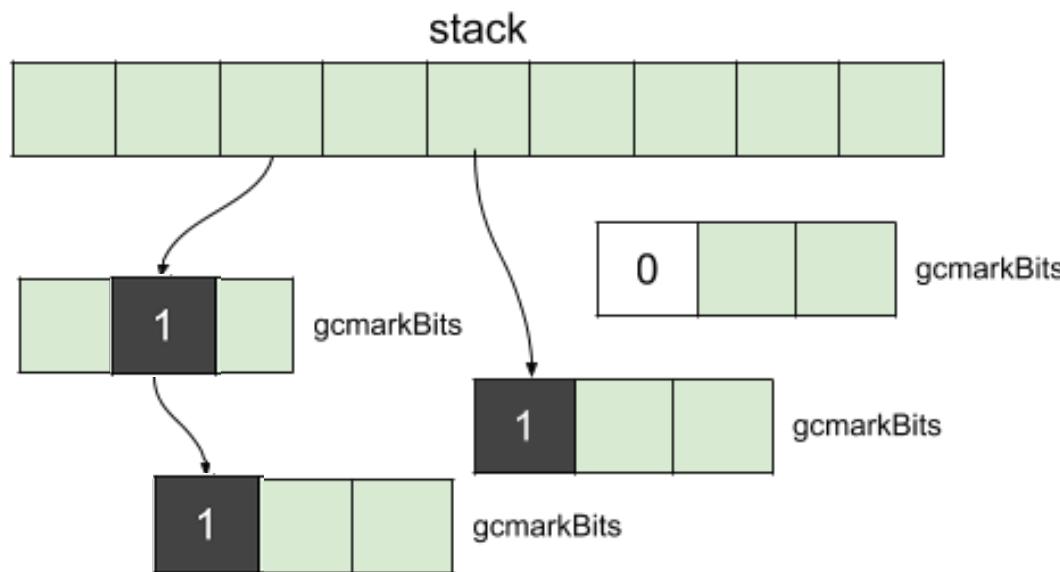
freelIndex

allocBits

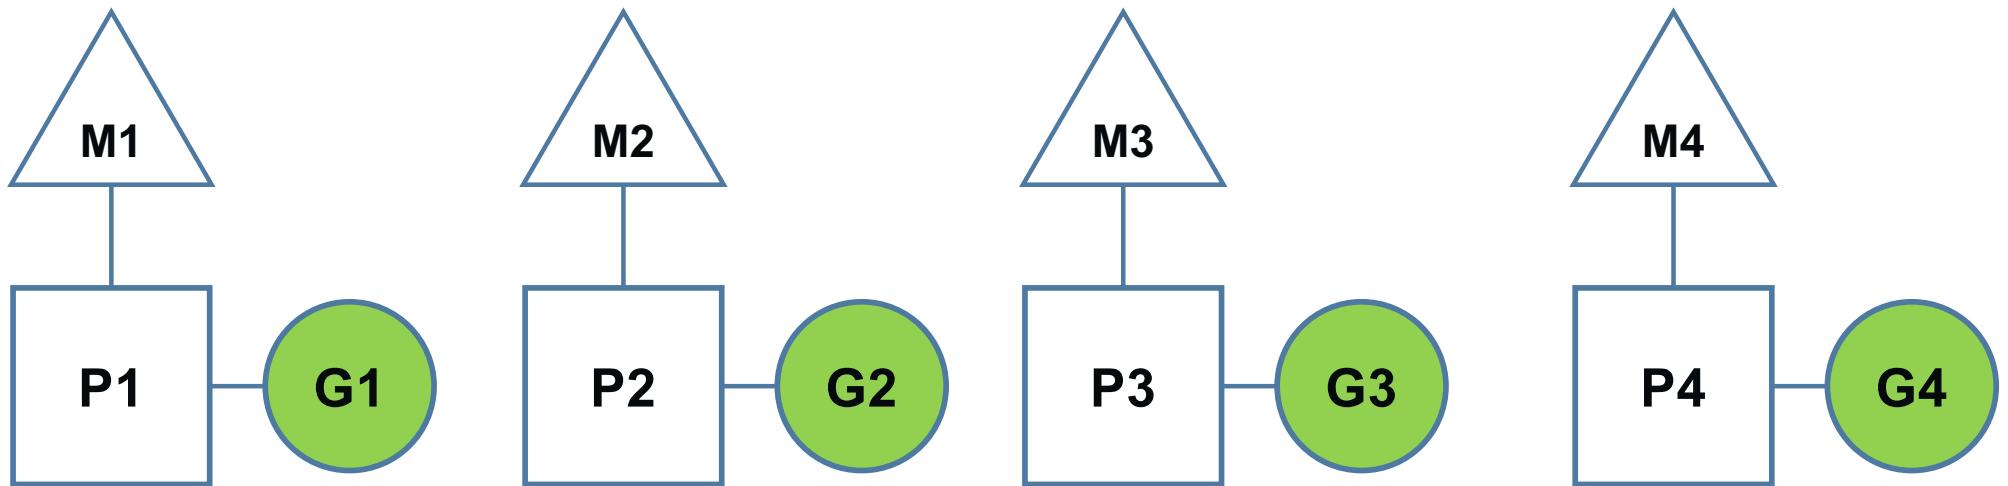


gcmarkBits

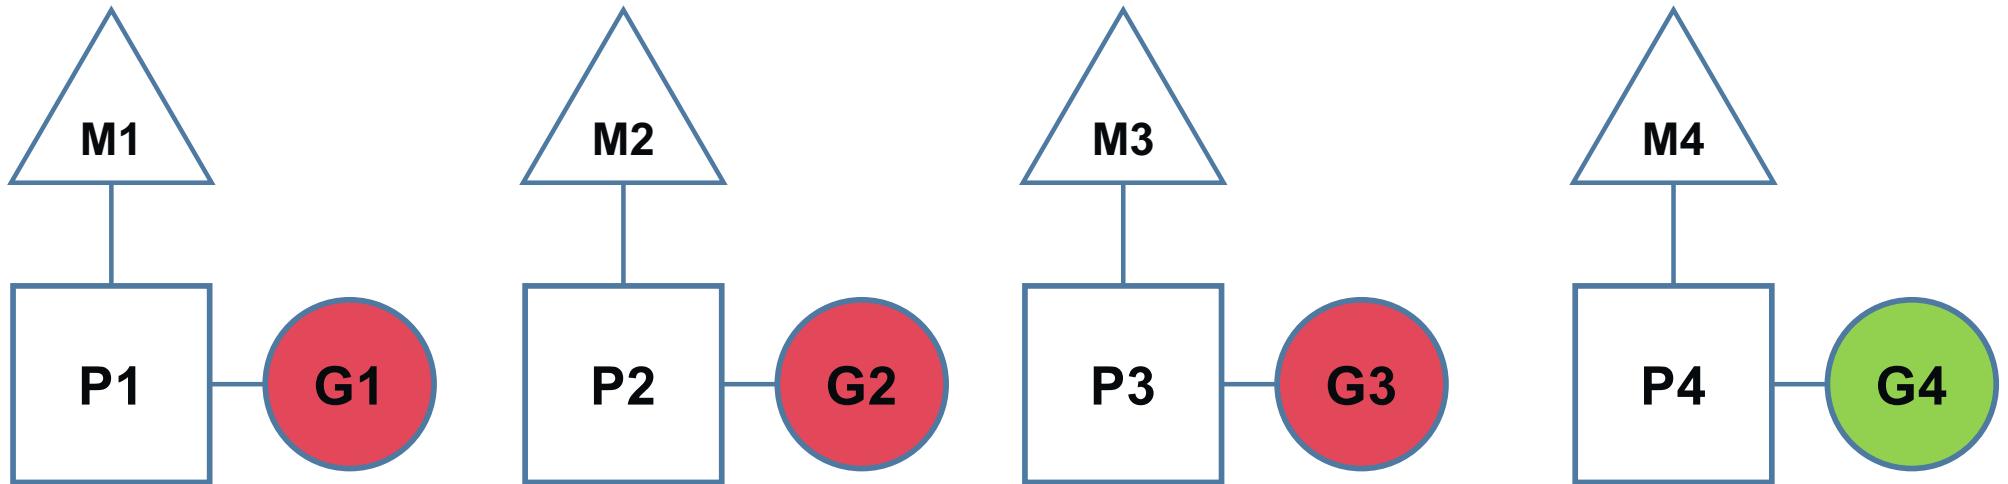




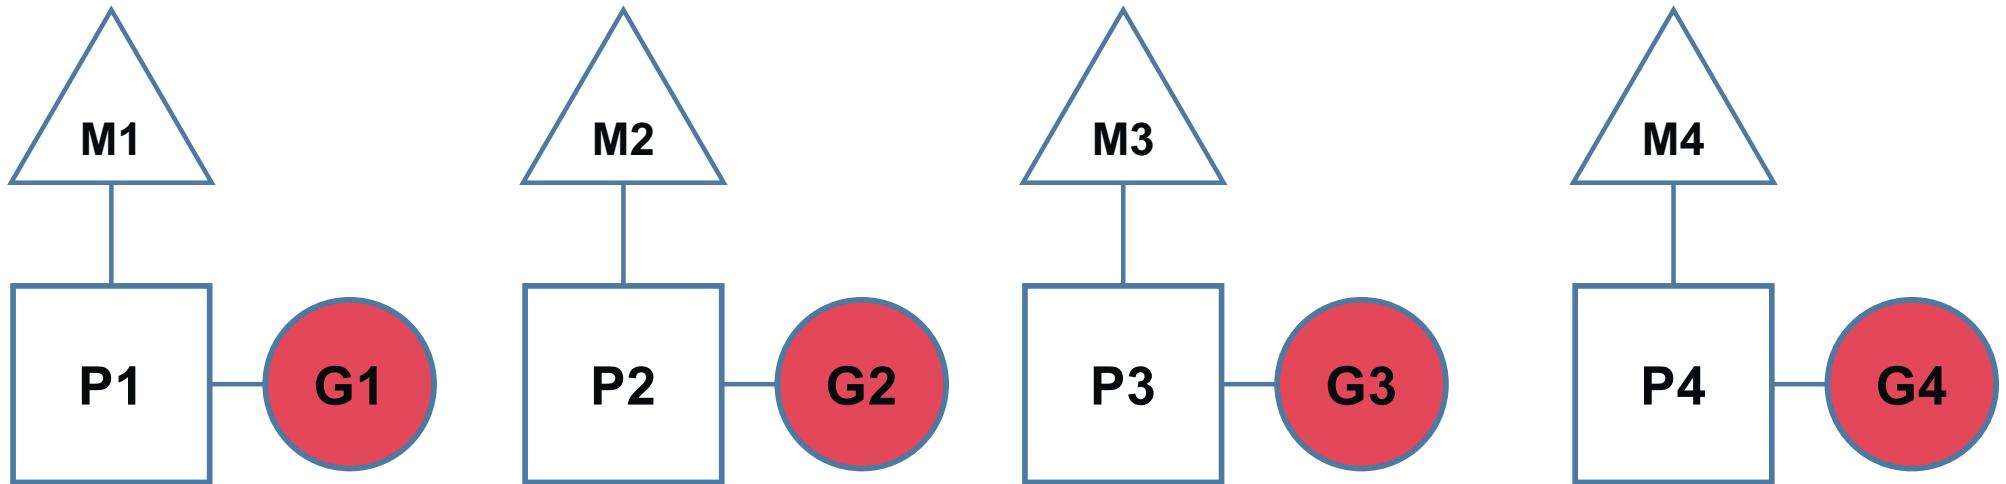
Mark Setup - STW



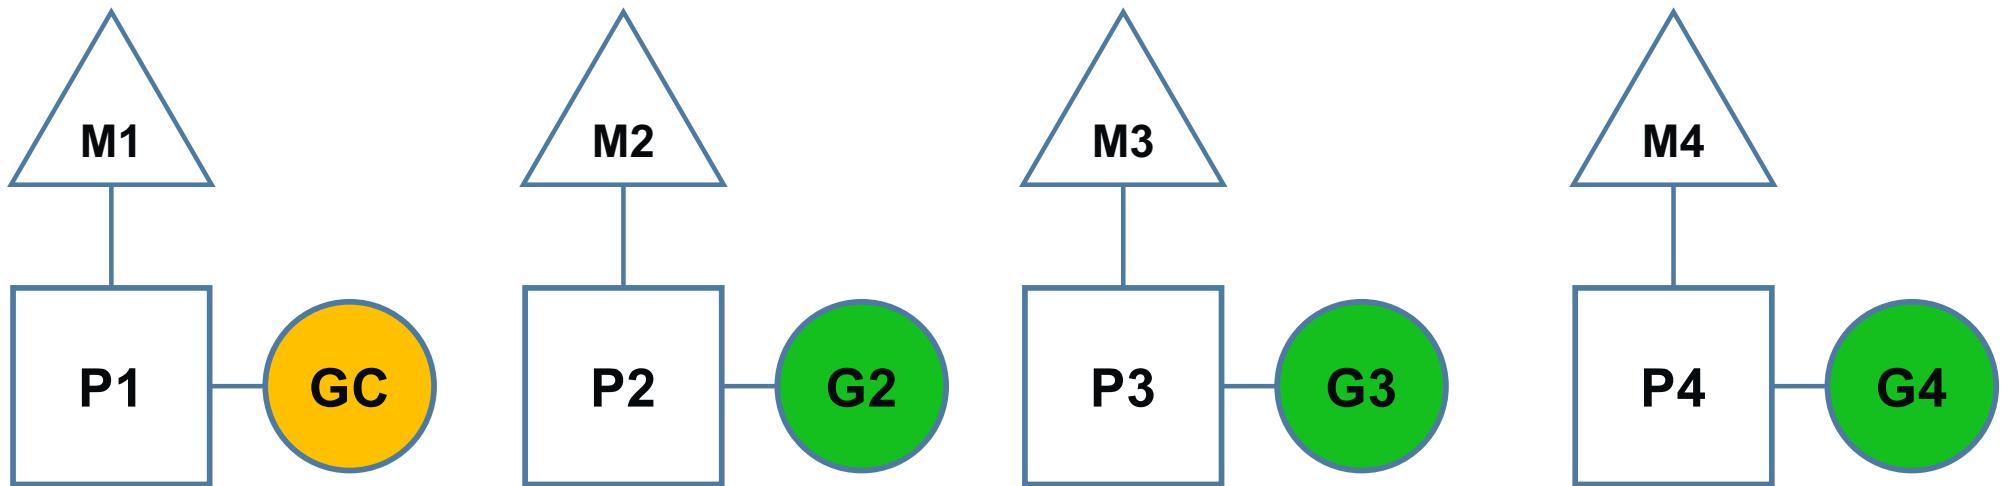
Mark Setup - STW



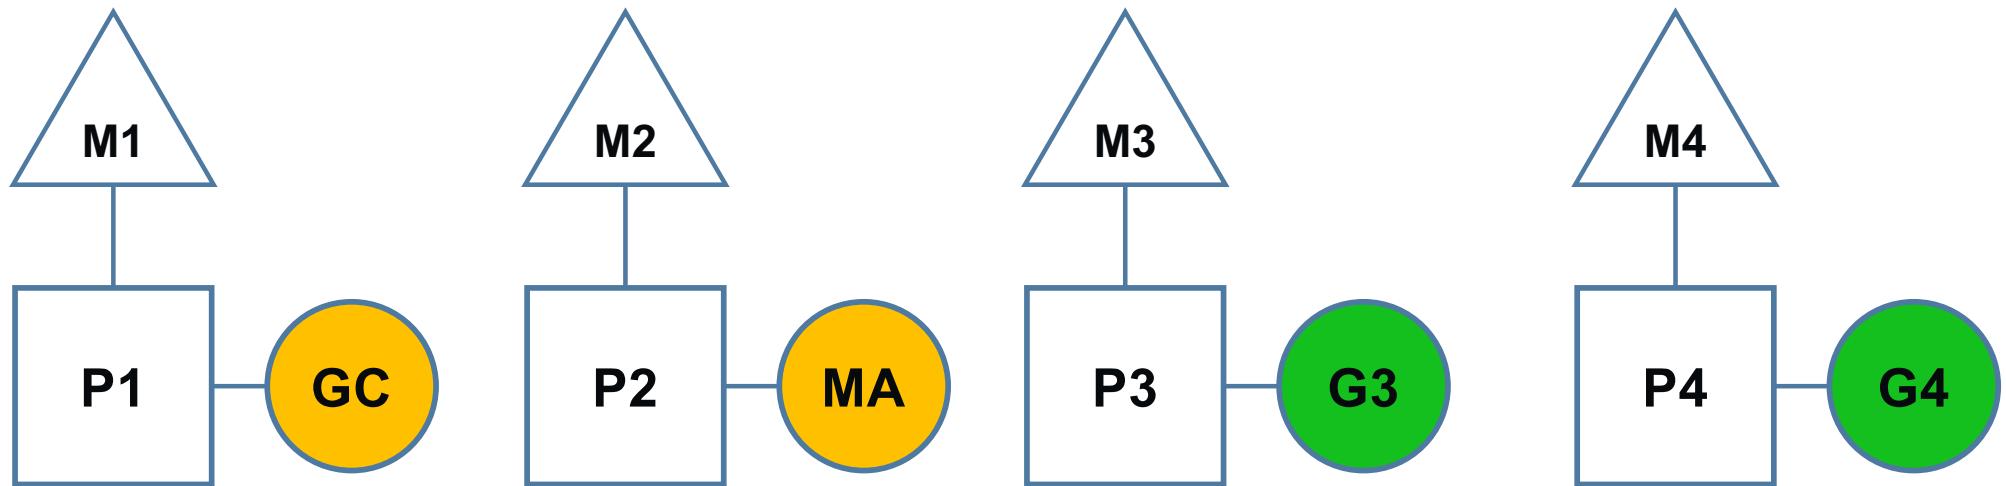
Mark Setup - STW



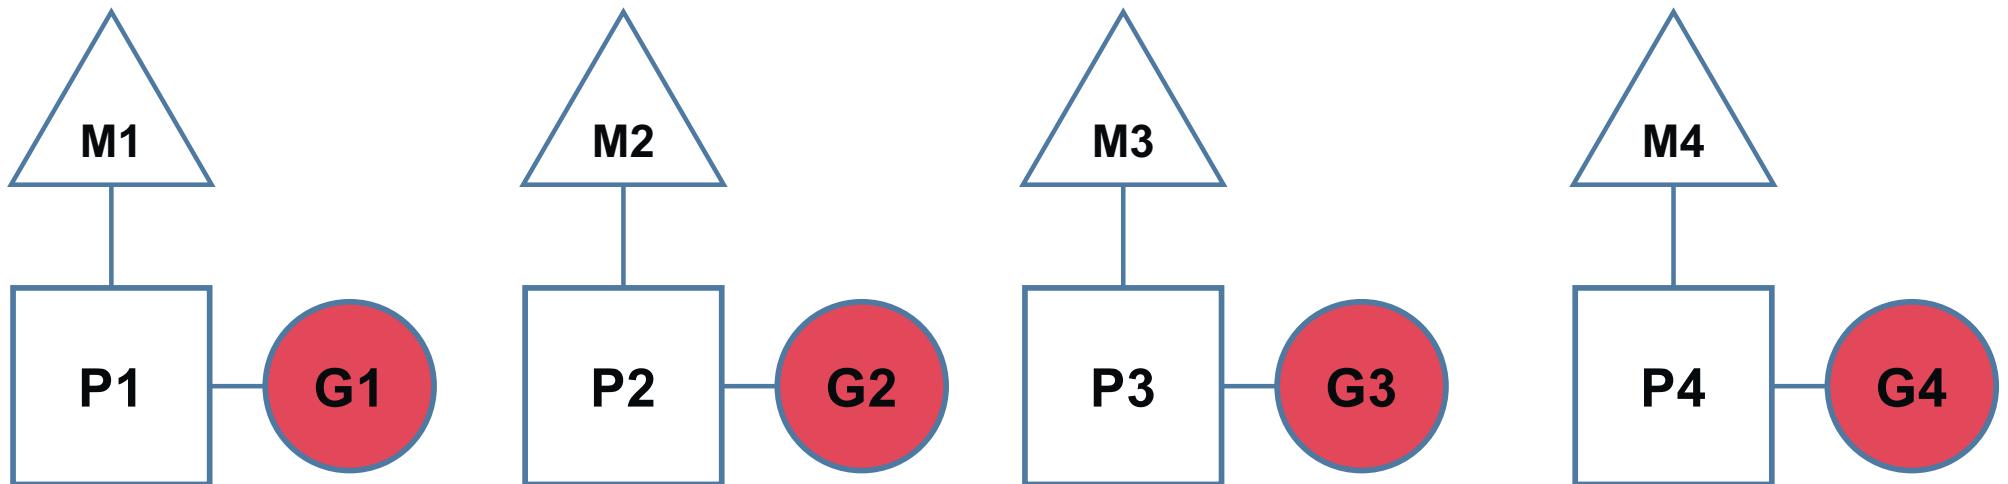
Marking - Concurrent

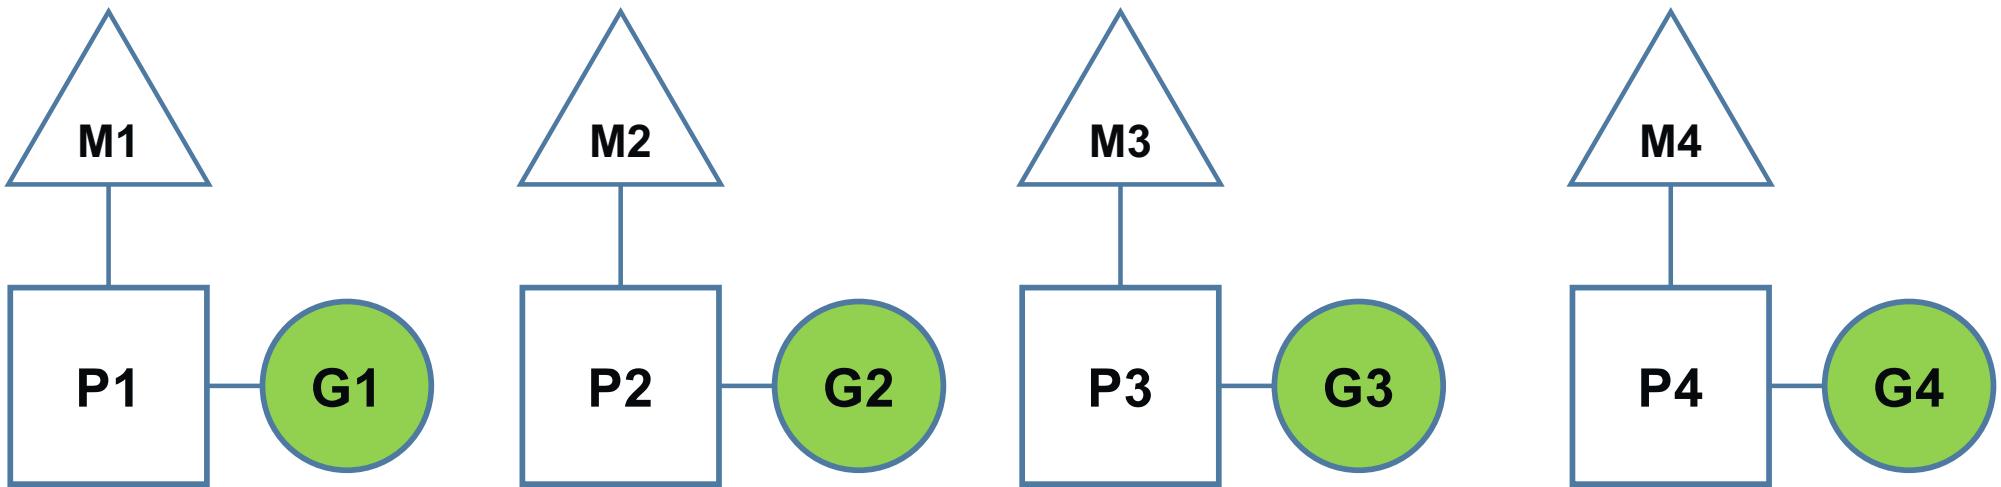


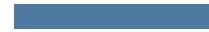
Mark Assist



Mark Termination - STW



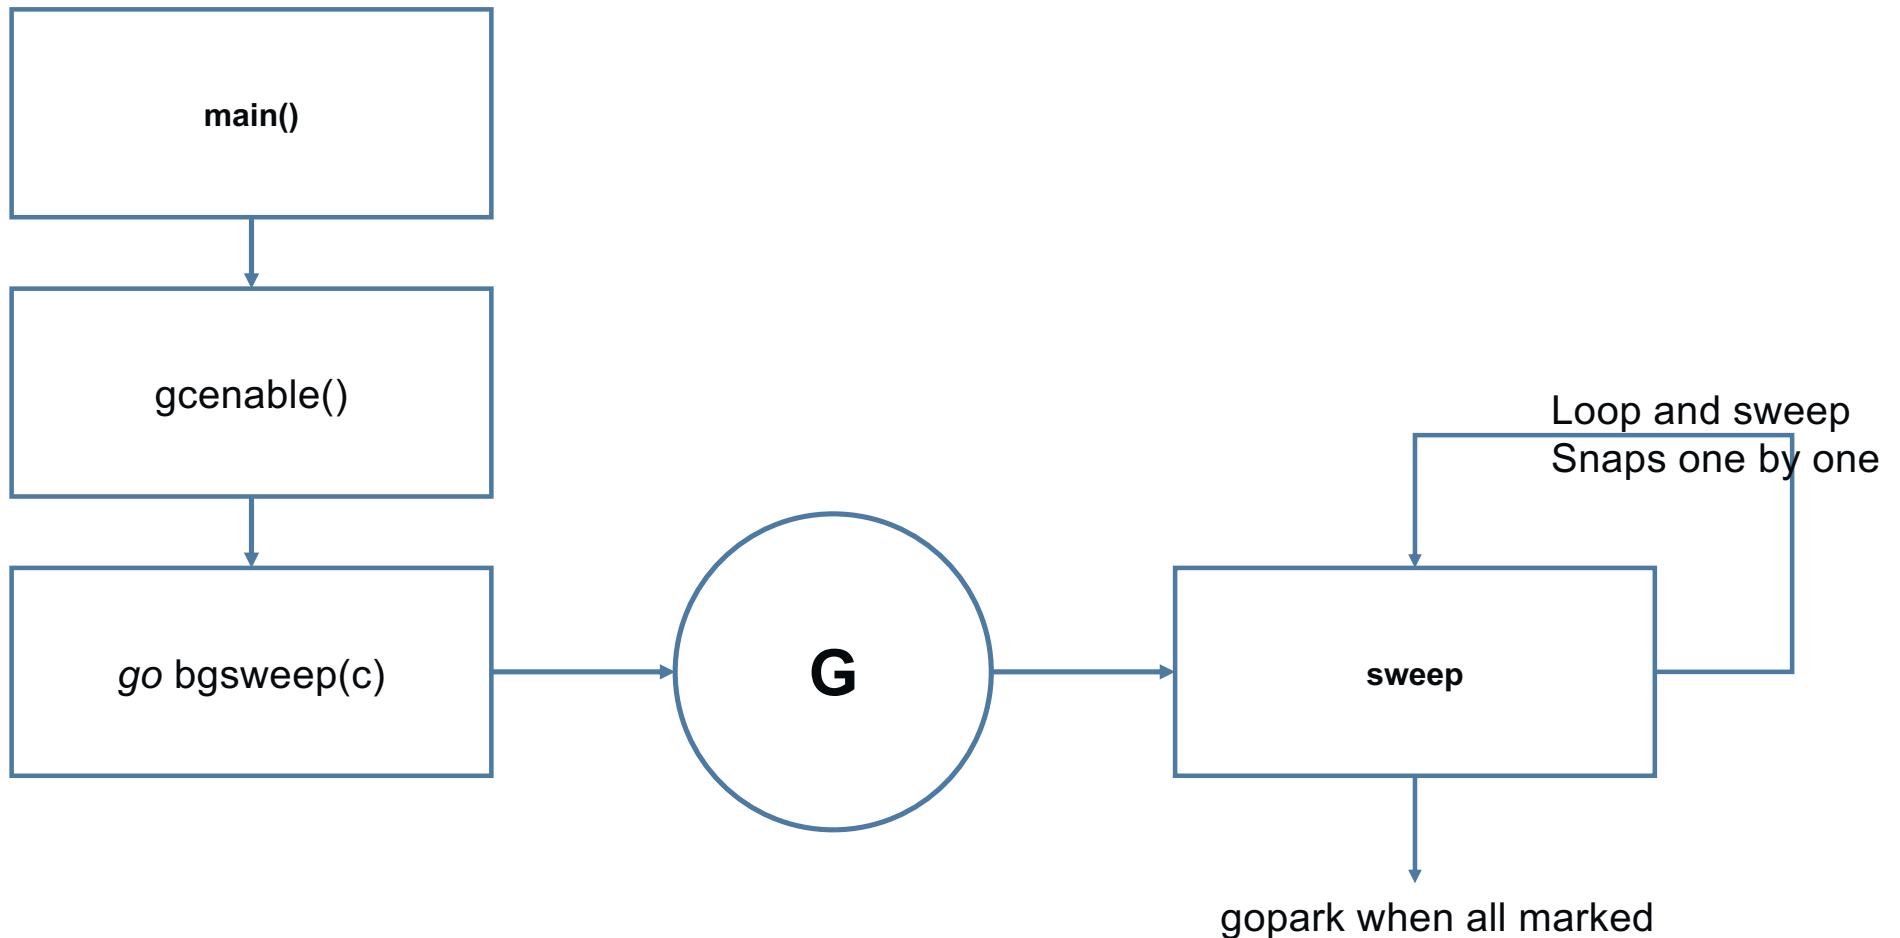




Sweeping - Concurrent

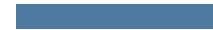
- Go provides two ways to sweep the memory:
- with a worker waiting in background that sweeps the spans one by one
- on the fly when the allocation requires a span

Sweep





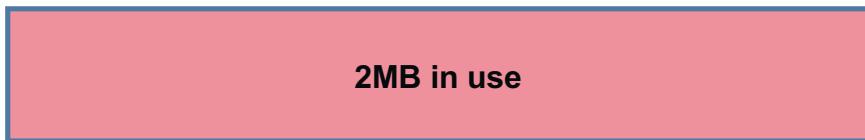
Collision with Collection Cycles



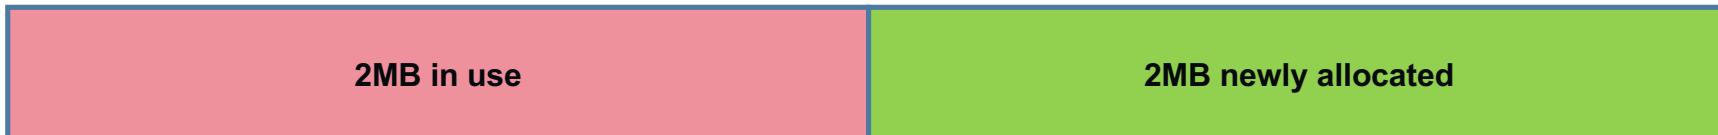
Timing

GC Percentage

$$\text{HeapGoal} = \text{HeapLive} * (1 + \text{GOGC}/100)$$



GOGC = 100





GC Trace

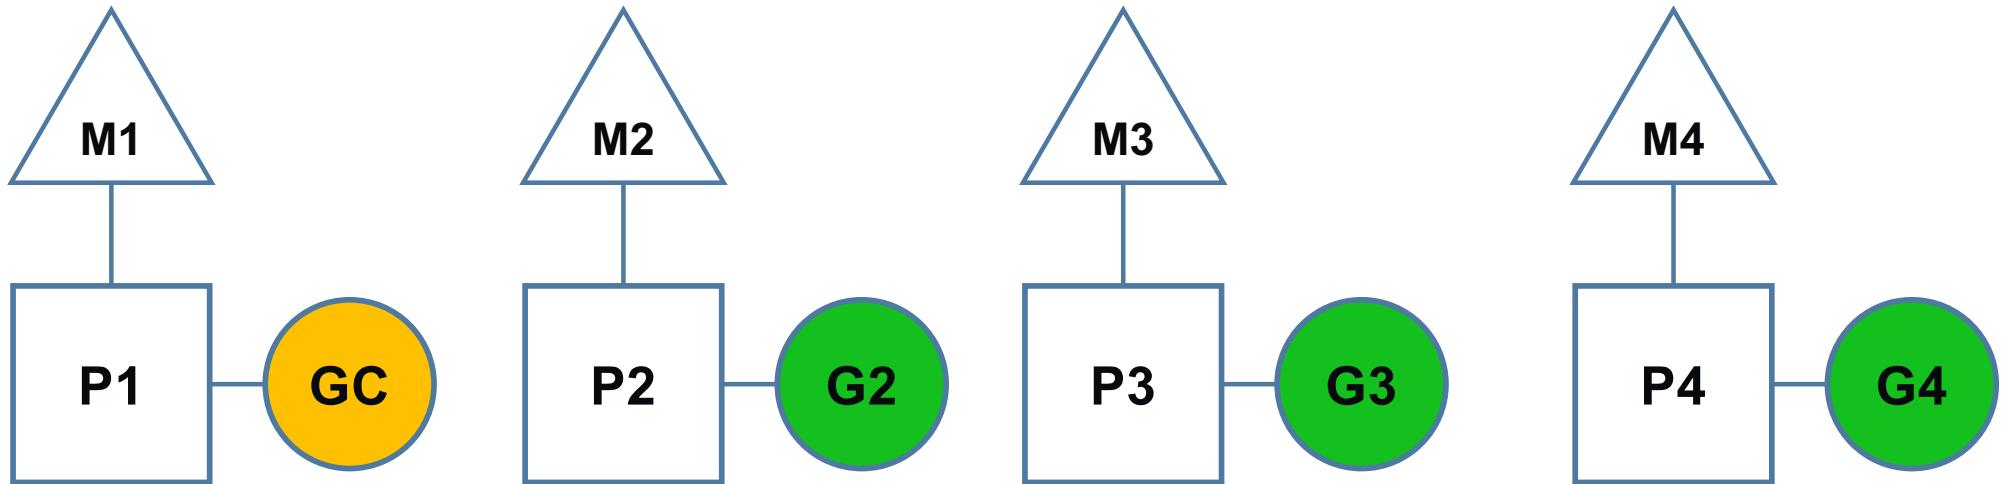
- GODEBUG=gctrace=1

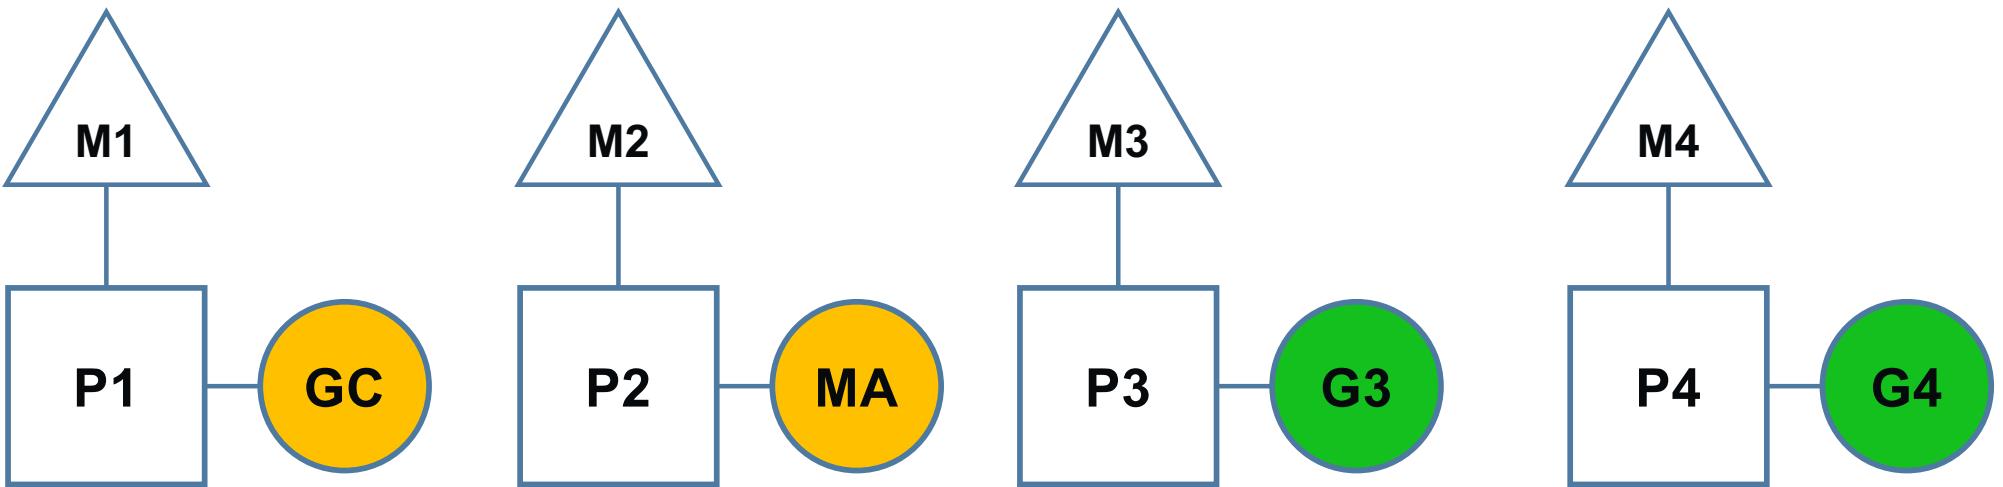
```
GODEBUG=gctrace=1 ./project > /dev/null
gc 3 @3.182s 0%: 0.015+0.59+0.096 ms clock, 0.19+0.10/1.3/3.0+1.1 ms cpu, 4->4->2 MB, 5 MB goal, 12 P
.
.
.
gc 2553 @8.452s 14%: 0.004+0.33+0.051 ms clock, 0.056+0.12/0.56/0.94+0.61 ms cpu, 4->4->2 MB, 5 MB goal, 12 P
```

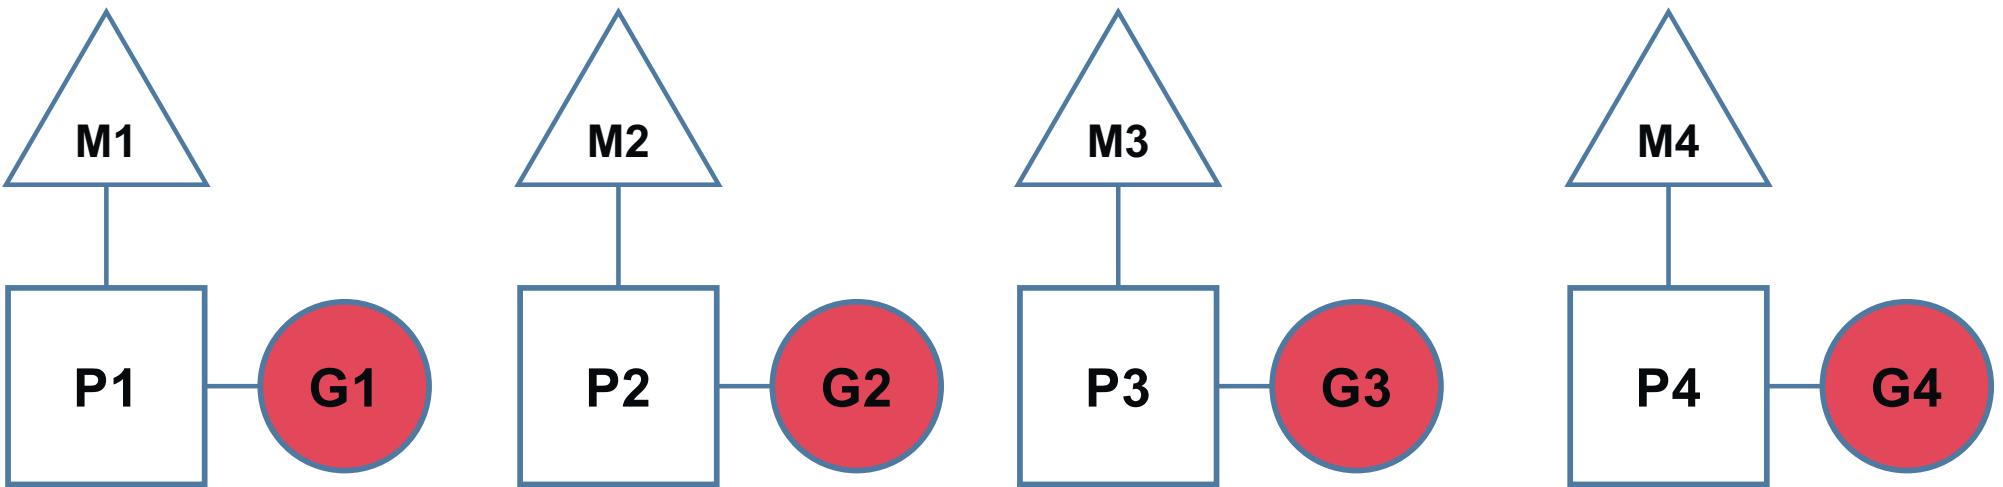


Pacing

Collector Latency Costs





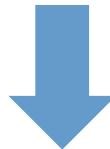




Less allocations on the heap



Less marking work



Shorter GC cycle

Values vs Pointers

- Prefer coping to passing pointers
- Avoid pointers in you types
- Check standard library for hidden pointers

```
// type Time struct {
//   wall and ext encode the wall time seconds, wall time nanoseconds,
//   and optional monotonic clock reading in nanoseconds.
//
//   From high to low bit position, wall encodes a 1-bit flag (hasMonotonic),
//   a 33-bit seconds field, and a 30-bit wall time nanoseconds field.
//   The nanoseconds field is in the range [0, 999999999].
//   If the hasMonotonic bit is 0, then the 33-bit field must be zero
//   and the full signed 64-bit wall seconds since Jan 1 year 1 is stored in ext.
//   If the hasMonotonic bit is 1, then the 33-bit field holds a 33-bit
//   unsigned wall seconds since Jan 1 year 1885, and ext holds a
//   signed 64-bit monotonic clock reading, nanoseconds since process start.
wall uint64
ext int64

// loc specifies the Location that should be used to
// determine the minute, hour, month, day, and year
// that correspond to this Time.
// The nil location means UTC.
// All UTC times are represented with loc=nil, never loc=&utcLoc.
loc *Location
}
```

Struct Layout

```
// Type size: 80 bytes
type S struct {
    a string
    b bool
    c string
    d *string
    e byte
    f []byte
}
```

Fields	Alignment
a string	8 bytes
b bool	1 byte
padding	8 bytes
c string	8 bytes
d pointer	8 bytes
e byte	1 byte
padding	8 bytes
f slice	8 bytes

```
// Type size: 72 bytes
type S struct {
    e byte
    b bool
    a string
    c string
    d *string
    f []byte
}
```

Fields	Alignment							
e byte								
b bool								
padding								
a string								
c string								
d pointer								
f slice								



Reuse Memory

- Reuse objects
- sync.Pool

Q&A

